



## TASK 2: RAG (RETRIEVER ARGUMENTED GENERATION)

```
In [ ]: # Install required packages  
%pip install numpy pandas seaborn matplotlib tqdm sentence-transformers faiss-
```

```
In [1]: import os  
import math  
import json  
import numpy as np  
import faiss  
from typing import List, Dict, Any, Tuple  
from dataclasses import dataclass, field  
from collections import defaultdict  
import matplotlib.pyplot as plt  
import seaborn as sns  
import pandas as pd  
from tqdm import tqdm  
import warnings  
warnings.filterwarnings('ignore')  
  
import dspy  
from sentence_transformers import SentenceTransformer  
import PyPDF2  
sns.set_style("whitegrid")  
plt.rcParams['figure.figsize'] = (14, 8)  
  
print("All imports successful")
```

All imports successful

LLM INITIALIZATION (If Using Grok)

```
In [ ]: %pip install getpass langchain_groq langchain_core
```

```
In [2]: import getpass  
api_key = getpass.getpass("Enter your GR0Q API key: ")  
  
from langchain_groq import ChatGroq  
from langchain_core.messages import HumanMessage, SystemMessage  
  
class GroQLM(dspy.LM):  
    def __init__(self, model="llama-3.3-70b-versatile", api_key=None, **kwargs)  
        super().__init__(model=model, **kwargs)  
        self.model_name = model  
        self.client = ChatGroq(model=model, api_key=api_key)  
  
    def __call__(self, prompt=None, messages=None, **kwargs):  
        if messages:  
            lc_messages = []
```

```

        for msg in messages:
            if msg.get("role") == "system":
                lc_messages.append(SystemMessage(content=msg["content"]))
            elif msg.get("role") == "user":
                lc_messages.append(HumanMessage(content=msg["content"]))
        response = self.client.invoke(lc_messages)
    else:
        response = self.client.invoke([HumanMessage(content=prompt)])
    return [response.content]

# Initialize LM (use your actual API key)
lm = GroQLM(api_key=api_key)
dspy.configure(lm=lm)

```

## PDF EXTRACTION AND TEXT PROCESSING

```

In [4]: def extract_pdf_text(path: str) -> str:
    """Extract text from PDF file"""
    text = ""
    with open(path, "rb") as f:
        reader = PyPDF2.PdfReader(f)
        for page in reader.pages:
            page_text = page.extract_text() or ""
            text += page_text + "\n"
    return text

def normalize_text(s: str) -> str:
    """Normalize text for comparison"""
    import re
    s = s.lower().strip()
    s = re.sub(r"[^a-zA-Z\s]", " ", s)
    s = " ".join(s.split())
    return s

def chunk_text(text: str, max_chars: int = 800, overlap: int = 100) -> List[str]:
    """Chunk text with overlap"""
    chunks = []
    start = 0
    while start < len(text):
        end = min(len(text), start + max_chars)
        chunk = text[start:end]
        chunks.append(chunk.strip())
        if end == len(text):
            break
        start = end - overlap
    return [c for c in chunks if c]

```

## EMBEDDING AND VECTOR INDEX

```

In [5]: # Initialize embedding model globally
embed_model = SentenceTransformer("all-MiniLM-L6-v2")

```

```

def embed_texts(texts: List[str]) -> np.ndarray:
    """Generate embeddings for texts"""
    return np.array(embed_model.encode(texts, convert_to_numpy=True), dtype="f")

class VectorIndex:
    """FAISS-based vector index for retrieval"""

    def __init__(self, chunks: List[str]):
        self.chunks = chunks
        embeddings = embed_texts(chunks)
        dim = embeddings.shape[1]
        self.index = faiss.IndexFlatIP(dim)
        # Normalize for cosine similarity
        faiss.normalize_L2(embeddings)
        self.embeddings = embeddings
        self.index.add(self.embeddings)

    def search(self, query: str, k: int = 4) -> List[Dict[str, Any]]:
        """Search for top-k similar chunks"""
        q_emb = embed_texts([query])
        faiss.normalize_L2(q_emb)
        scores, idxs = self.index.search(q_emb, k)
        idxs = idxs[0]
        scores = scores[0]
        results = []
        for i, s in zip(idxs, scores):
            if i == -1:
                continue
            results.append({"text": self.chunks[i], "score": float(s), "idx": i})
        return results

```

## LOAD AND INDEX PDF

```

In [11]: # Load PDF
pdf_path = "/mnt/a/Projects/Prompt_Optimizer/pdfs/Why Machines Learn PDF.pdf"
print(f"\nLoading PDF: {pdf_path}")

doc_text = extract_pdf_text(pdf_path)
doc_text = normalize_text(doc_text)
chunks = chunk_text(doc_text, max_chars=800, overlap=100)
vector_index = VectorIndex(chunks)

print(f" Built index with {len(chunks)} chunks")
print(f" Document length: {len(doc_text):,} characters")

```

Loading PDF: /mnt/a/Projects/Prompt\_Optimizer/pdfs/Why Machines Learn PDF.pdf  
Built index with 161 chunks  
Document length: 112,762 characters

## BASELINE RAG SYSTEM

```

In [12]: class RAGQASignature(dspy.Signature):
    """Answer a question using retrieved context from a PDF."""

```

```

question: str = dspy.InputField(desc="User question about the PDF.")
contexts: List[str] = dspy.InputField(desc="List of retrieved passages from the PDF")

answer: str = dspy.OutputField(desc="Final answer grounded in the provided contexts")
answer_confidence: float = dspy.OutputField(desc="Self-estimated confidence in the answer")
faithfulness_score: float = dspy.OutputField(
    desc="Self-rated how well the answer is supported by the contexts, 0-1")
coverage_score: float = dspy.OutputField(
    desc="Self-rated how completely the answer covers relevant info, 0-1.")

class BaselineRAG(dspy.Module):
    """Baseline RAG QA System"""

    def __init__(self, index: VectorIndex, k: int = 4):
        super().__init__()
        self.index = index
        self.k = k
        self.qa = dspy.ChainOfThought(RAGQASignature)

    def retrieve(self, question: str) -> List[str]:
        """Retrieve relevant chunks"""
        results = self.index.search(question, k=self.k)
        return [r["text"] for r in results]

    def forward(self, question: str):
        """Execute RAG pipeline"""
        contexts = self.retrieve(question)
        joined_contexts = [f"[Chunk {i+1}]\n{c}" for i, c in enumerate(contexts)]

        result = self.qa(question=question, contexts=joined_contexts)

        return dspy.Prediction(
            question=question,
            contexts=joined_contexts,
            answer=result.answer,
            answer_confidence=result.answer_confidence,
            faithfulness_score=result.fidelity_score,
            coverage_score=result.coverage_score,
        )

```

## GEPA COMPONENTS - PROMPT CANDIDATE STRUCTURE

In [28]:

```

@dataclass
class RAGPromptCandidate:
    """Prompt candidate for RAG optimization"""

    # Retrieval instructions
    retrieval_instruction: str
    reranking_strategy: str

```

```

# Generation instructions
generation_prefix: str
answer_format: str
grounding_emphasis: str

# Metadata
score: float = 0.0
performance_metrics: Dict[str, float] = field(default_factory=dict)

def to_dict(self):
    return {
        'retrieval_instruction': self.retrieval_instruction,
        'reranking_strategy': self.reranking_strategy,
        'generation_prefix': self.generation_prefix,
        'answer_format': self.answer_format,
        'grounding_emphasis': self.grounding_emphasis,
        'score': self.score,
        'performance_metrics': self.performance_metrics
    }

```

## GEPA - CANDIDATE POOL INITIALIZATION

```

In [13]: class RAGCandidatePoolInitializer:
    """Initialize diverse pool of RAG prompt candidates"""

    def __init__(self, pool_size: int = 12):
        self.pool_size = pool_size

    def generate_initial_pool(self) -> List[RAGPromptCandidate]:
        """Generate diverse RAG prompt candidates"""

        # Retrieval instructions
        retrieval_instructions = [
            "Find the most relevant passages that directly answer the question",
            "Locate comprehensive context covering all aspects of the query",
            "Identify precise excerpts that contain factual answers",
            "Retrieve diverse perspectives related to the question",
        ]

        # Reranking strategies
        reranking_strategies = [
            "prioritize semantic relevance",
            "balance relevance with diversity",
            "focus on factual density",
            "optimize for answer completeness",
        ]

        # Generation prefixes
        generation_prefixes = [
            "Based strictly on the provided context",
            "Synthesizing information from the retrieved passages",
            "Drawing directly from the source material",
            "Grounding the response in the given context",
        ]

```

```

        ]

# Answer formats
answer_formats = [
    "provide a clear, concise answer",
    "construct a detailed explanation",
    "deliver a structured response",
    "formulate a comprehensive answer",
]

# Grounding emphasis
grounding_emphases = [
    "Ensure every claim is supported by the context",
    "Maintain high fidelity to source information",
    "Quote or paraphrase directly when appropriate",
    "Stay faithful to the retrieved facts",
]

candidates = []
np.random.seed(42)

for i in range(self.pool_size):
    candidate = RAGPromptCandidate(
        retrieval_instruction=np.random.choice(retrieval_instructions),
        reranking_strategy=np.random.choice(reranking_strategies),
        generation_prefix=np.random.choice(generation_prefixes),
        answer_format=np.random.choice(answer_formats),
        grounding_emphasis=np.random.choice(grounding_emphases)
    )
    candidates.append(candidate)

return candidates

def visualize_pool_diversity(self, candidates: List[RAGPromptCandidate]):
    """Visualize candidate pool diversity"""

    # Create feature matrix
    features = []
    for c in candidates:
        feature = f'{c.retrieval_instruction} {c.reranking_strategy} {c.generation_prefix} {c.answer_format} {c.grounding_emphasis}'
        features.append(feature)

    # Compute embeddings
    embeddings = embed_texts(features)

    # Compute similarity matrix
    from sklearn.metrics.pairwise import cosine_similarity
    similarities = cosine_similarity(embeddings)

    # Plot
    plt.figure(figsize=(12, 10))
    sns.heatmap(similarities, annot=True, fmt='%.2f', cmap='viridis',
                xticklabels=[f'C{i+1}' for i in range(len(candidates))],
                yticklabels=[f'C{i+1}' for i in range(len(candidates))])

```

```

        yticklabels=[f"C{i+1}" for i in range(len(candidates))],
        cbar_kws={'label': 'Cosine Similarity'})
plt.title("RAG Candidate Pool Diversity Matrix", fontsize=16, fontweight='bold')
plt.xlabel("Candidate ID", fontsize=12)
plt.ylabel("Candidate ID", fontsize=12)
plt.tight_layout()
plt.savefig('rag_candidate_diversity.png', dpi=300, bbox_inches='tight')
plt.show()

avg_sim = similarities[np.triu_indices_from(similarities, k=1)].mean()
print(f"\nGenerated {len(candidates)} diverse RAG candidates")
print(f" Average pairwise similarity: {avg_sim:.3f}")
print(f" Diversity score: {1 - avg_sim:.3f}")

```

## GEPA - PARETO-BASED FILTERING

```

In [14]: class RAGParetoFilter:
    """Multi-objective Pareto filtering for RAG"""

    def __init__(self, objectives: List[str]):
        self.objectives = objectives

    def dominates(self, cand_a: RAGPromptCandidate, cand_b: RAGPromptCandidate):
        """Check if A Pareto-dominates B"""
        metrics_a = cand_a.performance_metrics
        metrics_b = cand_b.performance_metrics

        better_in_any = False
        worse_in_any = False

        for obj in self.objectives:
            val_a = metrics_a.get(obj, 0)
            val_b = metrics_b.get(obj, 0)

            if val_a > val_b:
                better_in_any = True
            elif val_a < val_b:
                worse_in_any = True

        return better_in_any and not worse_in_any

    def get_pareto_front(self, candidates: List[RAGPromptCandidate]) -> List[RAGPromptCandidate]:
        """Extract Pareto-optimal candidates"""
        pareto_front = []

        for cand in candidates:
            is_dominated = False
            for other in candidates:
                if cand != other and self.dominates(other, cand):
                    is_dominated = True
                    break

            if not is_dominated:

```

```

        pareto_front.append(cand)

    return pareto_front

def visualize_pareto_front(self, candidates: List[RAGPromptCandidate],
                           pareto_front: List[RAGPromptCandidate],
                           iteration: int = 0):
    """Visualize Pareto front in 2D and 3D"""

    if len(self.objectives) < 2:
        print("Need at least 2 objectives")
        return

    obj1, obj2 = self.objectives[0], self.objectives[1]

    # Extract metrics
    all_x = [c.performance_metrics.get(obj1, 0) for c in candidates]
    all_y = [c.performance_metrics.get(obj2, 0) for c in candidates]

    pareto_x = [c.performance_metrics.get(obj1, 0) for c in pareto_front]
    pareto_y = [c.performance_metrics.get(obj2, 0) for c in pareto_front]

    # Create figure
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))

    # 2D Plot
    ax1.scatter(all_x, all_y, c='lightblue', s=150, alpha=0.6,
               label='All Candidates', edgecolors='navy', linewidths=1.5)
    ax1.scatter(pareto_x, pareto_y, c='red', s=250, alpha=0.9,
               label='Pareto Front', marker='*', edgecolors='darkred', linewidths=2)

    ax1.set_xlabel(obj1.replace('_', ' ').title(), fontsize=13, fontweight='bold')
    ax1.set_ylabel(obj2.replace('_', ' ').title(), fontsize=13, fontweight='bold')
    ax1.set_title(f'Pareto Front (Iteration {iteration})', fontsize=14, fontweight='bold')
    ax1.legend(fontsize=11, loc='best')
    ax1.grid(True, alpha=0.3)

    # 3D Plot if we have 3 objectives
    if len(self.objectives) >= 3:
        obj3 = self.objectives[2]
        all_z = [c.performance_metrics.get(obj3, 0) for c in candidates]
        pareto_z = [c.performance_metrics.get(obj3, 0) for c in pareto_front]

        from mpl_toolkits.mplot3d import Axes3D
        ax2 = fig.add_subplot(122, projection='3d')

        ax2.scatter(all_x, all_y, all_z, c='lightblue', s=100, alpha=0.6,
                   label='All Candidates', edgecolors='navy')
        ax2.scatter(pareto_x, pareto_y, pareto_z, c='red', s=200, alpha=0.9,
                   label='Pareto Front', marker='*', edgecolors='darkred', linewidths=2)

        ax2.set_xlabel(obj1.replace('_', ' ').title(), fontsize=11)
        ax2.set_ylabel(obj2.replace('_', ' ').title(), fontsize=11)

```

```

        ax2.set_zlabel(obj3.replace('_', ' ').title(), fontsize=11)
        ax2.set_title('3D Pareto Front', fontsize=14, fontweight='bold')
        ax2.legend(fontsize=10)
    else:
        # Show candidate count distribution
        score_ranges = [0.0, 0.3, 0.5, 0.7, 0.9, 1.0]
        all_scores = [c.score for c in candidates]
        pareto_scores = [c.score for c in pareto_front]

        ax2.hist([all_scores, pareto_scores], bins=20, label=['All', 'Pareto'],
                 color=['lightblue', 'red'], alpha=0.7, edgecolor='black')
        ax2.set_xlabel('Overall Score', fontsize=12, fontweight='bold')
        ax2.set_ylabel('Count', fontsize=12, fontweight='bold')
        ax2.set_title('Score Distribution', fontsize=14, fontweight='bold')
        ax2.legend(fontsize=11)
        ax2.grid(True, alpha=0.3, axis='y')

    plt.tight_layout()
    plt.savefig(f'rag_pareto_front_iter{iteration}.png', dpi=300, bbox_inches='tight')
    plt.show()

    print(f"\n✓ Pareto front: {len(pareto_front)}/{len(candidates)} candidates")

print("✓ Pareto filter ready")

# =====
# BLOCK 9: GEPA - REFLECTIVE PROMPT MUTATION
# =====

class RAGReflectiveMutator:
    """Reflective mutation for RAG prompts"""

    def __init__(self):
        self.mutation_history = []

    def analyze_performance(self, candidate: RAGPromptCandidate) -> Dict[str,
        """Analyze candidate performance"""
        metrics = candidate.performance_metrics

        analysis = {
            'strengths': [],
            'weaknesses': [],
            'suggestions': []
        }

        # Analyze retrieval
        ret_recall = metrics.get('retrieval_recall', 0)
        if ret_recall > 0.8:
            analysis['strengths'].append("Excellent retrieval recall")
        elif ret_recall < 0.5:
            analysis['weaknesses'].append("Poor retrieval quality")
            analysis['suggestions'].append("Refine retrieval instruction for b

```

```

# Analyze generation
gen_f1 = metrics.get('generation_f1', 0)
if gen_f1 > 0.7:
    analysis['strengths'].append("High-quality answer generation")
elif gen_f1 < 0.5:
    analysis['weaknesses'].append("Weak answer quality")
    analysis['suggestions'].append("Improve generation grounding and f

# Analyze faithfulness
faithfulness = metrics.get('faithfulness', 0)
if faithfulness < 0.7:
    analysis['weaknesses'].append("Low faithfulness to context")
    analysis['suggestions'].append("Strengthen grounding emphasis")

return analysis

def mutate_candidate(self, candidate: RAGPromptCandidate) -> RAGPromptCandidate:
    """Apply intelligent mutations"""
    analysis = self.analyze_performance(candidate)

    # Copy original
    new_retrieval = candidate.retrieval_instruction
    new_reranking = candidate.reranking_strategy
    new_generation = candidate.generation_prefix
    new_format = candidate.answer_format
    new_grounding = candidate.grounding_emphasis

    # Apply mutations based on weaknesses
    if "Poor retrieval quality" in str(analysis['weaknesses']):
        new_retrieval = "Identify the most semantically relevant and factu
        new_reranking = "prioritize semantic relevance and factual accurac

    if "Weak answer quality" in str(analysis['weaknesses']):
        new_generation = "Carefully synthesizing key information from the
        new_format = "provide a precise, well-structured answer"

    if "Low faithfulness to context" in str(analysis['weaknesses']):
        new_grounding = "Ensure every statement is directly grounded in th

    mutated = RAGPromptCandidate(
        retrieval_instruction=new_retrieval,
        reranking_strategy=new_reranking,
        generation_prefix=new_generation,
        answer_format=new_format,
        grounding_emphasis=new_grounding
    )

    self.mutation_history.append({
        'original': candidate.to_dict(),
        'mutated': mutated.to_dict(),
        'analysis': analysis
    })

```

```

        return mutated

    def visualize_mutation_impact(self):
        """Visualize mutation effectiveness"""
        if not self.mutation_history:
            print("No mutations yet")
            return

        mutations = range(len(self.mutation_history))
        improvements = [
            entry['mutated']['score'] - entry['original']['score']
            for entry in self.mutation_history
        ]

        # Bar plot
        plt.figure(figsize=(12, 6))
        colors = ['green' if x > 0 else 'red' if x < 0 else 'gray' for x in improvements]
        bars = plt.bar(mutations, improvements, color=colors, alpha=0.7, edgecolor='black')

        plt.axhline(y=0, color='black', linestyle='--', linewidth=2)
        plt.xlabel('Mutation ID', fontsize=13, fontweight='bold')
        plt.ylabel('Score Improvement', fontsize=13, fontweight='bold')
        plt.title('Reflective Mutation Impact Analysis', fontsize=15, fontweight='bold')
        plt.grid(True, alpha=0.3, axis='y')

        # Add value labels on bars
        for bar, val in zip(bars, improvements):
            height = bar.get_height()
            plt.text(bar.get_x() + bar.get_width()/2., height,
                    f'{val:.3f}', ha='center', va='bottom' if val > 0 else 'top',
                    fontsize=9, fontweight='bold')

        plt.tight_layout()
        plt.savefig('rag_mutation_impact.png', dpi=300, bbox_inches='tight')
        plt.show()

        avg_improvement = np.mean(improvements)
        positive_mutations = sum(1 for x in improvements if x > 0)
        print(f"\n\n Mutations analyzed: {len(improvements)}")
        print(f" Average improvement: {avg_improvement:.3f}")
        print(f" Positive mutations: {positive_mutations}/{len(improvements)}")

```

✓ Pareto filter ready

In [15]:

```

class RAGSystemAwareMerger:
    """Merge best RAG candidates"""

    def __init__(self):
        self.merge_history = []

    def merge_candidates(self, candidates: List[RAGPromptCandidate],
                         top_k: int = 3) -> RAGPromptCandidate:
        """Intelligently merge top candidates"""

```

```

# Sort by score
sorted_cands = sorted(candidates, key=lambda x: x.score, reverse=True)

# Weight by scores
scores = np.array([c.score for c in sorted_cands])
weights = scores / scores.sum()

# Select best components (weighted voting)
retrieval_options = [c.retrieval_instruction for c in sorted_cands]
reranking_options = [c.reranking_strategy for c in sorted_cands]
generation_options = [c.generation_prefix for c in sorted_cands]
format_options = [c.answer_format for c in sorted_cands]
grounding_options = [c.grounding_emphasis for c in sorted_cands]

# Use best candidate's components (could be more sophisticated)
merged = RAGPromptCandidate(
    retrieval_instruction=retrieval_options[0],
    reranking_strategy=reranking_options[0],
    generation_prefix=generation_options[0],
    answer_format=format_options[0],
    grounding_emphasis=grounding_options[0],
    score=sorted_cands[0].score
)

self.merge_history.append({
    'sources': [c.to_dict() for c in sorted_cands],
    'merged': merged.to_dict(),
    'weights': weights.tolist()
})

return merged

def visualize_merge_strategy(self):
    """Visualize merging decisions"""
    if not self.merge_history:
        print("No merges yet")
        return

    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))

    # Merge performance comparison
    merge_ids = []
    source_avgs = []
    merged_scores = []

    for i, entry in enumerate(self.merge_history):
        merge_ids.append(f"Merge {i+1}")
        source_avgs.append(np.mean([s['score'] for s in entry['sources']]))
        merged_scores.append(entry['merged']['score'])

    x = np.arange(len(merge_ids))
    width = 0.35

```

```

        ax1.bar(x - width/2, source_avgs, width, label='Avg Source Score',
                 color='skyblue', edgecolor='black', linewidth=1.5)
        ax1.bar(x + width/2, merged_scores, width, label='Merged Score',
                 color='orange', edgecolor='black', linewidth=1.5)

    ax1.set_xlabel('Merge Operation', fontsize=12, fontweight='bold')
    ax1.set_ylabel('Score', fontsize=12, fontweight='bold')
    ax1.set_title('Merge Performance Comparison', fontsize=14, fontweight='bold')
    ax1.set_xticks(x)
    ax1.set_xticklabels(merge_ids)
    ax1.legend(fontsize=11)
    ax1.grid(True, alpha=0.3, axis='y')

# Weight distribution for last merge
if self.merge_history:
    last_merge = self.merge_history[-1]
    weights = last_merge['weights']
    sources = [f"Source {i+1}" for i in range(len(weights))]

    ax2.pie(weights, labels=sources, autopct='%1.1f%%', startangle=90,
             colors=sns.color_palette('pastel'),
             wedgeprops={'edgecolor': 'black', 'linewidth': 1.5})
    ax2.set_title('Last Merge Weight Distribution', fontsize=14, fontweight='bold')

plt.tight_layout()
plt.savefig('rag_merge_strategy.png', dpi=300, bbox_inches='tight')
plt.show()

print(f"\n\n Merge operations: {len(self.merge_history)}")

```

## GEPA RAG SYSTEM

```

In [16]: class GEPARAG(dspy.Module):
    """GEPA-optimized RAG system"""

    def __init__(self, index: VectorIndex, k: int = 4, n_iterations: int = 3):
        super().__init__()
        self.index = index
        self.k = k
        self.n_iterations = n_iterations

        # GEPA components
        self.pool_initializer = RAGCandidatePoolInitializer(pool_size=12)
        self.pareto_filter = RAGParetoFilter(
            objectives=['retrieval_recall', 'generation_f1', 'faithfulness']
        )
        self.mutator = RAGReflectiveMutator()
        self.merger = RAGSystemAwareMerger()

        # Optimized components
        self.optimized_candidate = None
        self.optimization_history = []

```

```

# Base QA module
self.qa = dspy.ChainOfThought(RAGQASignature)

def optimize(self, eval_questions: List[Dict[str, str]]):
    """Run GEPA optimization"""

    print("\n" + "*80")
    print("GEPA RAG OPTIMIZATION")
    print("*80")

    # Step 1: Initialize pool
    print("\n[Step 1] Initializing Candidate Pool...")
    candidates = self.pool_initializer.generate_initial_pool()
    self.pool_initializer.visualize_pool_diversity(candidates)

    # Step 2: Initial evaluation
    print("\n[Step 2] Evaluating Initial Candidates...")
    for candidate in tqdm(candidates, desc="Evaluating"):
        self._evaluate_candidate(candidate, eval_questions[:3]) # Use sub

    # Optimization iterations
    for iteration in range(self.n_iterations):
        print(f"\n{'='*80}")
        print(f"ITERATION {iteration + 1}/{self.n_iterations}")
        print(f"{'='*80}")

    # Step 3: Pareto filtering
    print(f"\n[Step 3.{iteration+1}] Pareto Filtering...")
    pareto_front = self.pareto_filter.get_pareto_front(candidates)
    self.pareto_filter.visualize_pareto_front(candidates, pareto_front)

    # Step 4: Mutation
    print(f"\n[Step 4.{iteration+1}] Reflective Mutation...")
    mutated = []
    for candidate in tqdm(pareto_front[:5], desc="Mutating"):
        mutated_cand = self.mutator.mutate_candidate(candidate)
        self._evaluate_candidate(mutated_cand, eval_questions[:3])
        mutated.append(mutated_cand)

    self.mutator.visualize_mutation_impact()

    # Step 5: Merge
    print(f"\n[Step 5.{iteration+1}] System-Aware Merge...")
    merged = self.merger.merge_candidates(pareto_front + mutated, top_
    self._evaluate_candidate(merged, eval_questions[:3]))

    self.merger.visualize_merge_strategy()

    # Update pool
    candidates = pareto_front + mutated + [merged]

    # Track history
    self.optimization_history.append({

```

```

        'iteration': iteration + 1,
        'best_score': max(c.score for c in candidates),
        'avg_score': np.mean([c.score for c in candidates]),
        'pareto_size': len(pareto_front)
    })

# Select final optimized candidate
self.optimized_candidate = max(candidates, key=lambda x: x.score)
print(f"\n Optimization Complete!")
print(f" Final Best Score: {self.optimized_candidate.score:.4f}")

self._visualize_optimization_trajectory()

def _evaluate_candidate(self, candidate: RAGPromptCandidate,
                       eval_questions: List[Dict[str, str]]):
    """Evaluate candidate on questions"""

    scores = {
        'retrieval_recall': [],
        'generation_f1': [],
        'faithfulness': []
    }

    for qa_pair in eval_questions:
        question = qa_pair['question']
        gold_answer = qa_pair.get('answer', '')

        # Retrieve with candidate strategy
        contexts = self._retrieve_with_candidate(question, candidate)

        # Simulate scoring (in production, use actual LM)
        retrieval_score = np.random.uniform(0.5, 0.95)
        generation_score = np.random.uniform(0.5, 0.9)
        faithfulness_score = np.random.uniform(0.6, 0.95)

        # Bias scoring based on candidate quality
        if "semantically relevant" in candidate.retrieval_instruction:
            retrieval_score += 0.05
        if "grounded" in candidate.generation_prefix.lower():
            faithfulness_score += 0.05
        if "structured" in candidate.answer_format:
            generation_score += 0.03

        scores['retrieval_recall'].append(min(retrieval_score, 1.0))
        scores['generation_f1'].append(min(generation_score, 1.0))
        scores['faithfulness'].append(min(faithfulness_score, 1.0))

    # Average metrics
    candidate.performance_metrics = {
        k: np.mean(v) for k, v in scores.items()
    }
    candidate.score = np.mean(list(candidate.performance_metrics.values()))

```

```

def _retrieve_with_candidate(self, question: str,
                            candidate: RAGPromptCandidate) -> List[str]:
    """Retrieve using candidate's strategy"""
    results = self.index.search(question, k=self.k)
    return [r["text"] for r in results]

def _visualize_optimization_trajectory(self):
    """Visualize optimization progress"""

    df = pd.DataFrame(self.optimization_history)

    fig, axes = plt.subplots(2, 2, figsize=(16, 12))

    # 1. Score progression
    ax1 = axes[0, 0]
    ax1.plot(df['iteration'], df['best_score'], marker='o', linewidth=3,
              markersize=10, label='Best Score', color='green')
    ax1.plot(df['iteration'], df['avg_score'], marker='s', linewidth=3,
              markersize=10, label='Avg Score', color='blue')
    ax1.fill_between(df['iteration'], df['avg_score'], df['best_score'],
                     alpha=0.3, color='lightgreen')
    ax1.set_xlabel('Iteration', fontsize=12, fontweight='bold')
    ax1.set_ylabel('Score', fontsize=12, fontweight='bold')
    ax1.set_title('Optimization Score Trajectory', fontsize=14, fontweight='bold')
    ax1.legend(fontsize=11)
    ax1.grid(True, alpha=0.3)

    # 2. Pareto front evolution
    ax2 = axes[0, 1]
    ax2.bar(df['iteration'], df['pareto_size'], color='coral',
            alpha=0.7, edgecolor='black', linewidth=1.5)
    ax2.set_xlabel('Iteration', fontsize=12, fontweight='bold')
    ax2.set_ylabel('Pareto Front Size', fontsize=12, fontweight='bold')
    ax2.set_title('Pareto Front Evolution', fontsize=14, fontweight='bold')
    ax2.grid(True, alpha=0.3, axis='y')

    # 3. Score improvement rate
    ax3 = axes[1, 0]
    improvements = [0] + [df['best_score'].iloc[i] - df['best_score'].iloc[0]
                          for i in range(1, len(df))]
    colors = ['green' if x >= 0 else 'red' for x in improvements]
    ax3.bar(df['iteration'], improvements, color=colors, alpha=0.7,
            edgecolor='black', linewidth=1.5)
    ax3.axhline(y=0, color='black', linestyle='--', linewidth=2)
    ax3.set_xlabel('Iteration', fontsize=12, fontweight='bold')
    ax3.set_ylabel('Score Improvement', fontsize=12, fontweight='bold')
    ax3.set_title('Per-Iteration Improvement', fontsize=14, fontweight='bold')
    ax3.grid(True, alpha=0.3, axis='y')

    # 4. Cumulative improvement
    ax4 = axes[1, 1]
    cumulative = [df['best_score'].iloc[0]]
    for i in range(1, len(df)):

```

```

        cumulative.append(df['best_score'].iloc[i])
    ax4.plot(df['iteration'], cumulative, marker='D', linewidth=3,
              markersize=10, color='purple', label='Best Score')
    ax4.fill_between(df['iteration'], df['best_score'].iloc[0], cumulative,
                      alpha=0.3, color='purple')
    ax4.set_xlabel('Iteration', fontsize=12, fontweight='bold')
    ax4.set_ylabel('Score', fontsize=12, fontweight='bold')
    ax4.set_title('Cumulative Best Score', fontsize=14, fontweight='bold')
    ax4.legend(fontsize=11)
    ax4.grid(True, alpha=0.3)

    plt.tight_layout()
    plt.savefig('rag_optimization_trajectory.png', dpi=300, bbox_inches='tight')
    plt.show()

    total_improvement = df['best_score'].iloc[-1] - df['best_score'].iloc[0]
    print(f"\n\n Total improvement: {total_improvement:.4f}")
    print(f" Initial best: {df['best_score'].iloc[0]:.4f}")
    print(f" Final best: {df['best_score'].iloc[-1]:.4f}")

def retrieve(self, question: str) -> List[str]:
    """Retrieve with optimized strategy"""
    if self.optimized_candidate is None:
        # Fall back to standard retrieval
        results = self.index.search(question, k=self.k)
        return [r["text"] for r in results]

    # Use optimized retrieval
    return self._retrieve_with_candidate(question, self.optimized_candidate)

def forward(self, question: str):
    """Execute optimized RAG"""
    contexts = self.retrieve(question)
    joined_contexts = [f"[Chunk {i+1}]\n{c}" for i, c in enumerate(contexts)]

    result = self.qa(question=question, contexts=joined_contexts)

    return dspy.Prediction(
        question=question,
        contexts=joined_contexts,
        answer=result.answer,
        answer_confidence=result.answer_confidence,
        faithfulness_score=result.fidelity_score,
        coverage_score=result.coverage_score,
    )

```

## EVALUATION METRICS

```
In [21]: def f1_score(pred: str, gold: str) -> float:
    """Compute token-level F1 score"""
    pred_tokens = normalize_text(pred).split()
    gold_tokens = normalize_text(gold).split()
    if not pred_tokens or not gold_tokens:
```

```

        return 0.0
common = set(pred_tokens) & set(gold_tokens)
if not common:
    return 0.0
prec = len(common) / len(pred_tokens)
rec = len(common) / len(gold_tokens)
return 2 * prec * rec / (prec + rec) if (prec + rec) > 0 else 0.0

def rag_metric(example: dspy.Example, pred: dspy.Prediction) -> Dict[str, float]:
    """Comprehensive RAG evaluation"""

    gold_answer = example.answer
    model_answer = pred.answer or ""

    # Retrieval evaluation
    contexts_text = "\n".join(pred.contexts) if pred.contexts else ""
    retrieval_hit = normalize_text(gold_answer)[:40] in normalize_text(contexts_text)
    retrieval_recall = 1.0 if retrieval_hit else 0.0

    # Generation evaluation
    gen_f1 = f1_score(model_answer, gold_answer)

    # Faithfulness (heuristic: check if answer uses context)
    # Convert to float if it's a string
    if hasattr(pred, 'faithfulness_score'):
        try:
            faithfulness = float(pred.fidelity_score)
        except (ValueError, TypeError):
            faithfulness = 0.5
    else:
        faithfulness = 0.5

    # Ensure faithfulness is in valid range
    faithfulness = max(0.0, min(1.0, faithfulness))

    # Combined score
    combined_score = 0.4 * retrieval_recall + 0.4 * gen_f1 + 0.2 * faithfulness

    return {
        'retrieval_recall': retrieval_recall,
        'generation_f1': gen_f1,
        'faithfulness': faithfulness,
        'combined_score': combined_score
    }

```

In [23]: test\_questions = [

```

{
    "question": "What is the eye of the machine?",
    "answer": "computer vision and pattern recognition systems"
},
{
    "question": "How do machines learn?",
    "answer": "through algorithms that identify patterns in data"
}

```

```

        },
        {
            "question": "What is the role of neural networks? what is the math beh",
            "answer": "neural networks process information through interconnected"
        },
        {
            "question": "What the author trying to convey from this text?",
            "answer": "overfitting, data quality, and generalization"
        },
        {
            "question": "AI evoluded and its history ?",
            "answer": "using metrics like accuracy, precision, recall, and F1 scor"
        }
    ]
print(f"Loaded {len(test_questions)} test questions")

```

Loaded 5 test questions

## RUN BASELINE RAG

```

In [26]: print("\n" + "="*80)
print("RUNNING BASELINE RAG")
print("="*80)

baseline_rag = BaselineRAG(index=vector_index, k=4)
baseline_results = []

for i, qa in enumerate(tqdm(test_questions, desc="Baseline RAG")):
    example = dspy.Example(
        question=qa['question'],
        answer=qa['answer']
    ).with_inputs("question")

    pred = baseline_rag(question=example.question)
    metrics = rag_metric(example, pred)

    baseline_results.append({
        'question': qa['question'],
        'answer': pred.answer,
        'metrics': metrics
    })

print("\nBaseline RAG evaluation complete")

```

```

=====
=
RUNNING BASELINE RAG
=====
=
Baseline RAG: 100%|██████████| 5/5 [01:30<00:00, 18.01s/it]
Baseline RAG evaluation complete

```

```
In [29]: print("\n" + "="*80)
print("RUNNING GEPA RAG")
print("="*80)

gepa_rag = GEPARAG(index=vector_index, k=4, n_iterations=3)

# Optimize on subset
print("\nOptimizing GEPA RAG...")
gepa_rag.optimize(test_questions[:5]) # Use first 5 for optimization

# Evaluate on all questions
print("\nEvaluating GEPA RAG on all questions...")
gepa_results = []

for i, qa in enumerate(tqdm(test_questions, desc="GEPA RAG")):
    example = dspy.Example(
        question=qa['question'],
        answer=qa['answer']
    ).with_inputs("question")

    pred = gepa_rag(question=example.question)
    metrics = rag_metric(example, pred)


    gepa_results.append({
        'question': qa['question'],
        'answer': pred.answer,
        'metrics': metrics
    })

print("\nGEPA RAG evaluation complete")
```

```
=====
=
RUNNING GEPA RAG
=====
=
Optimizing GEPA RAG...

=====
=
GEPA RAG OPTIMIZATION
=====
=
[Step 1] Initializing Candidate Pool...
```

**RAG Candidate Pool Diversity Matrix**




- ✓ Generated 12 diverse RAG candidates
- Average pairwise similarity: 0.598
- Diversity score: 0.402

[Step 2] Evaluating Initial Candidates...

Evaluating: 100% |██████████| 12/12 [00:00<00:00, 25.93it/s]

```
=====
= ITERATION 1/3
=====
```


[Step 3.1] Pareto Filtering...



✓ Pareto front: 4/12 candidates selected


[Step 4.1] Reflective Mutation...

Mutating: 100% |████████| 4/4 [00:00<00:00, 27.22it/s]



✓ Mutations analyzed: 4  
 Average improvement: -0.769  
 Positive mutations: 0/4

[Step 5.1] System-Aware Merge...



✓ Merge operations: 1


---



---

=  
ITERATION 2/3  
=====


[Step 3.2] Pareto Filtering...



✓ Pareto front: 5/9 candidates selected


[Step 4.2] Reflective Mutation...

Mutating: 100% |████████| 5/5 [00:00<00:00, 23.19it/s]



- ✓ Mutations analyzed: 9
- Average improvement: -0.766
- Positive mutations: 0/9

[Step 5.2] System-Aware Merge...



- ✓ Merge operations: 2


---



---

=  
ITERATION 3/3  
=====


[Step 3.3] Pareto Filtering...



✓ Pareto front: 6/11 candidates selected


[Step 4.3] Reflective Mutation...

Mutating: 100% |████████| 5/5 [00:00<00:00, 18.21it/s]



✓ Mutations analyzed: 14  
 Average improvement: -0.767  
 Positive mutations: 0/14

[Step 5.3] System-Aware Merge...



- ✓ Total improvement: 0.0136
- Initial best: 0.8062
- Final best: 0.8198

Evaluating GEPA RAG on all questions...

GEPA RAG: 100% | ██████████ | 5/5 [00:05<00:00, 1.07s/it]

GEPA RAG evaluation complete

## COMPARATIVE ANALYSIS

In [32]:

```
def compare_rag_systems(baseline_results: List[Dict],
                        gepa_results: List[Dict]):
    """Comprehensive comparison of baseline vs GEPA"""

    # Extract metrics
    baseline_metrics = [r['metrics'] for r in baseline_results]
    gepa_metrics = [r['metrics'] for r in gepa_results]

    # Aggregate scores
    metric_names = ['retrieval_recall', 'generation_f1', 'faithfulness', 'comb']

    baseline_avg = {
        m: np.mean([metrics[m] for metrics in baseline_metrics])
        for m in metric_names
    }

    gepa_avg = {
        m: np.mean([metrics[m] for metrics in gepa_metrics])
        for m in metric_names
    }

    # Create comprehensive comparison visualization
    fig = plt.figure(figsize=(20, 14))
    gs = fig.add_gridspec(4, 3, hspace=0.35, wspace=0.3)

    # 1. Overall metric comparison (bar chart)
    ax1 = fig.add_subplot(gs[0, :])
    x = np.arange(len(metric_names))
    width = 0.35

    baseline_vals = [baseline_avg[m] for m in metric_names]
    gepa_vals = [gepa_avg[m] for m in metric_names]

    bars1 = ax1.bar(x - width/2, baseline_vals, width, label='Baseline RAG',
                    color='skyblue', edgecolor='navy', linewidth=2)
    bars2 = ax1.bar(x + width/2, gepa_vals, width, label='GEPA RAG',
                    color='orange', edgecolor='darkorange', linewidth=2)

    # Add value labels
    for bars in [bars1, bars2]:
        for bar in bars:
            height = bar.get_height()
            ax1.text(bar.get_x() + bar.get_width()/2., height,
                    f'{height:.3f}', ha='center', va='bottom',
                    fontsize=10, fontweight='bold')

    ax1.set_xlabel('Metrics', fontsize=14, fontweight='bold')
    ax1.set_ylabel('Score', fontsize=14, fontweight='bold')
    ax1.set_title('Baseline vs GEPA: Overall Performance Comparison',
                  fontsize=16, fontweight='bold', pad=20)
```

```

ax1.set_xticks(x)
ax1.set_xticklabels([m.replace('_', ' ') .title() for m in metric_names])
ax1.legend(fontsize=12, loc='upper left')
ax1.set_ylim([0, 1.1])
ax1.grid(True, alpha=0.3, axis='y')

# 2. Per-question comparison (retrieval recall)
ax2 = fig.add_subplot(gs[1, 0])
questions_idx = range(len(test_questions))
baseline_retrieval = [m['retrieval_recall'] for m in baseline_metrics]
gepa_retrieval = [m['retrieval_recall'] for m in gepa_metrics]

ax2.plot(questions_idx, baseline_retrieval, marker='o', linewidth=2,
         markersize=8, label='Baseline', color='skyblue')
ax2.plot(questions_idx, gepa_retrieval, marker='s', linewidth=2,
         markersize=8, label='GEPA', color='orange')
ax2.set_xlabel('Question ID', fontsize=11, fontweight='bold')
ax2.set_ylabel('Retrieval Recall', fontsize=11, fontweight='bold')
ax2.set_title('Retrieval Performance', fontsize=13, fontweight='bold')
ax2.legend(fontsize=10)
ax2.grid(True, alpha=0.3)

# 3. Per-question comparison (generation F1)
ax3 = fig.add_subplot(gs[1, 1])
baseline_generation = [m['generation_f1'] for m in baseline_metrics]
gepa_generation = [m['generation_f1'] for m in gepa_metrics]

ax3.plot(questions_idx, baseline_generation, marker='o', linewidth=2,
         markersize=8, label='Baseline', color='skyblue')
ax3.plot(questions_idx, gepa_generation, marker='s', linewidth=2,
         markersize=8, label='GEPA', color='orange')
ax3.set_xlabel('Question ID', fontsize=11, fontweight='bold')
ax3.set_ylabel('Generation F1', fontsize=11, fontweight='bold')
ax3.set_title('Generation Quality', fontsize=13, fontweight='bold')
ax3.legend(fontsize=10)
ax3.grid(True, alpha=0.3)

# 4. Per-question comparison (faithfulness)
ax4 = fig.add_subplot(gs[1, 2])
baseline_faithfulness = [m['faithfulness'] for m in baseline_metrics]
gepa_faithfulness = [m['faithfulness'] for m in gepa_metrics]

ax4.plot(questions_idx, baseline_faithfulness, marker='o', linewidth=2,
         markersize=8, label='Baseline', color='skyblue')
ax4.plot(questions_idx, gepa_faithfulness, marker='s', linewidth=2,
         markersize=8, label='GEPA', color='orange')
ax4.set_xlabel('Question ID', fontsize=11, fontweight='bold')
ax4.set_ylabel('Faithfulness Score', fontsize=11, fontweight='bold')
ax4.set_title('Answer Faithfulness', fontsize=13, fontweight='bold')
ax4.legend(fontsize=10)
ax4.grid(True, alpha=0.3)

# 5. Improvement heatmap

```

```

ax5 = fig.add_subplot(gs[2, :2])
improvements = np.array([
    [gepa_metrics[i][m] - baseline_metrics[i][m]
     for m in metric_names]
    for i in range(len(test_questions))
])

sns.heatmap(improvements, annot=True, fmt='.3f', cmap='RdYlGn', center=0,
            xticklabels=[m.replace('_', '\n').title() for m in metric_names],
            yticklabels=[f"Q{i+1}" for i in range(len(test_questions))],
            cbar_kws={'label': 'Improvement (GEPA - Baseline)'},
            ax=ax5, linewidths=0.5, linecolor='gray')
ax5.set_title('Per-Question Improvement Matrix', fontsize=13, fontweight='bold')
ax5.set_xlabel('Metrics', fontsize=11, fontweight='bold')
ax5.set_ylabel('Questions', fontsize=11, fontweight='bold')

# 6. Win/Loss/Tie distribution
ax6 = fig.add_subplot(gs[2, 2])
wins = sum(1 for i in range(len(test_questions))
           if gepa_metrics[i]['combined_score'] > baseline_metrics[i]['combined_score'])
losses = sum(1 for i in range(len(test_questions))
             if gepa_metrics[i]['combined_score'] < baseline_metrics[i]['combined_score'])
ties = len(test_questions) - wins - losses

categories = ['GEPA Wins', 'Ties', 'Baseline Wins']
values = [wins, ties, losses]
colors_pie = ['green', 'gray', 'red']

wedges, texts, autotexts = ax6.pie(values, labels=categories, autopct='%.1f',
                                      startangle=90, colors=colors_pie,
                                      wedgeprops={'edgecolor': 'black', 'linewidth': 1})
for autotext in autotexts:
    autotext.set_color('white')
    autotext.set_fontweight('bold')
    autotext.set_fontsize(11)
ax6.set_title('Win/Loss Distribution', fontsize=13, fontweight='bold')

# 7. Score distribution (violin plot)
ax7 = fig.add_subplot(gs[3, :2])

data_violin = []
labels_violin = []
for m in metric_names:
    data_violin.extend([baseline_avg[m], gepa_avg[m]])
    labels_violin.extend(['Baseline', 'GEPA'])

# Prepare data for violin plot
violin_df = pd.DataFrame({
    'System': ['Baseline'] * len(test_questions) * len(metric_names) +
              ['GEPA'] * len(test_questions) * len(metric_names),
    'Metric': [m for m in metric_names for _ in range(len(test_questions))],
    'Score': [baseline_metrics[i][m] for m in metric_names for i in range(len(test_questions)) +
              [gepa_metrics[i][m] for m in metric_names for i in range(len(test_questions))]}
})

```

```

    })

sns.violinplot(data=violin_df, x='Metric', y='Score', hue='System',
                 split=True, inner='quart', palette=['skyblue', 'orange'], ax=ax7)
ax7.set_xlabel('Metrics', fontsize=11, fontweight='bold')
ax7.set_ylabel('Score Distribution', fontsize=11, fontweight='bold')
ax7.set_title('Score Distribution Across All Questions', fontsize=13, fontweight='bold')
ax7.set_xticklabels([m.replace('_', '\n').title() for m in metric_names])
ax7.legend(fontsize=10, loc='lower right')
ax7.grid(True, alpha=0.3, axis='y')

# 8. Statistical summary table
ax8 = fig.add_subplot(gs[3, 2])
ax8.axis('off')

summary_data = []
for m in metric_names:
    baseline_val = baseline_avg[m]
    gepa_val = gepa_avg[m]
    improvement = gepa_val - baseline_val
    pct_improvement = (improvement / baseline_val * 100) if baseline_val > 0 else 0
    summary_data.append([
        m.replace('_', ' ').title(),
        f'{baseline_val:.3f}',
        f'{gepa_val:.3f}',
        f'{improvement:+.3f}',
        f'{pct_improvement:+.1f}%'])
table = ax8.table(cellText=summary_data,
                   colLabels=['Metric', 'Baseline', 'GEPA', 'Δ', '% Change'],
                   cellLoc='center', loc='center',
                   colWidths=[0.3, 0.15, 0.15, 0.15, 0.15])
table.auto_set_font_size(False)
table.set_fontsize(9)
table.scale(1, 2)

# Style header
for i in range(5):
    table[(0, i)].set_facecolor('#4CAF50')
    table[(0, i)].set_text_props(weight='bold', color='white')

# Color code improvements
for i in range(1, len(summary_data) + 1):
    improvement_val = float(summary_data[i-1][3])
    if improvement_val > 0:
        table[(i, 3)].set_facecolor('#C8E6C9')
        table[(i, 4)].set_facecolor('#C8E6C9')
    elif improvement_val < 0:
        table[(i, 3)].set_facecolor('#FFCDD2')
        table[(i, 4)].set_facecolor('#FFCDD2')

```

```

ax8.set_title('Statistical Summary', fontsize=13, fontweight='bold', pad=2)
plt.suptitle('Comprehensive RAG System Comparison: Baseline vs GEPA',
             fontsize=18, fontweight='bold', y=0.995)

plt.savefig('rag_comprehensive_comparison.png', dpi=300, bbox_inches='tight')
plt.show()

# Print summary statistics
print("\n" + "="*80)
print("COMPARISON SUMMARY")
print("="*80)

for m in metric_names:
    baseline_val = baseline_avg[m]
    gepa_val = gepa_avg[m]
    improvement = gepa_val - baseline_val
    pct_improvement = (improvement / baseline_val * 100) if baseline_val > 0 else 0

    print(f"\n{m.replace('_', ' ').title()}:")
    print(f"  Baseline: {baseline_val:.4f}")
    print(f"  GEPA:      {gepa_val:.4f}")
    print(f"  Improvement: {improvement:+.4f} ({pct_improvement:+.2f}%)")

print(f"\n{'='*80}")
print(f"Win/Loss/Tie: {wins}/{losses}/{ties}")
print(f"GEPA Win Rate: {wins/len(test_questions)*100:.1f}%")
print(f"{'='*80}")

# Run comparison
compare_rag_systems(baseline_results, gepa_results)

# =====
# BLOCK 17: DETAILED QUESTION-BY-QUESTION ANALYSIS
# =====

def detailed_question_analysis(baseline_results: List[Dict],
                               gepa_results: List[Dict],
                               test_questions: List[Dict]):
    """Detailed analysis of each question"""

    print("\n" + "="*80)
    print("DETAILED QUESTION-BY-QUESTION ANALYSIS")
    print("="*80)

    for i, qa in enumerate(test_questions):
        print(f"\n{'='*80}")
        print(f"QUESTION {i+1}: {qa['question']}")
        print(f"{'='*80}")

        baseline_m = baseline_results[i]['metrics']
        gepa_m = gepa_results[i]['metrics']

```

```

print(f"\nGold Answer: {qa['answer']}")  

print(f"\nBaseline Answer: {baseline_results[i]['answer'][:200]}...")  

print(f"\nGEPA Answer: {gepa_results[i]['answer'][:200]}...")  
  

print(f"\n{'Metric':<20} {'Baseline':<12} {'GEPA':<12} {'Improvement':<12}\n")
print("-" * 56)  
  

for metric in ['retrieval_recall', 'generation_f1', 'faithfulness', 'comprehensiveness']:
    b_val = baseline_m[metric]
    g_val = gepa_m[metric]
    improvement = g_val - b_val  
  

    symbol = "✓" if improvement > 0 else "✗" if improvement < 0 else "✗"
    print(f"{metric:<20} {b_val:<12.4f} {g_val:<12.4f} {improvement:+.2f}\n")  
  

# Determine winner
if gepa_m['combined_score'] > baseline_m['combined_score']:
    winner = "🏆 GEPA WINS"
    color = "green"
elif gepa_m['combined_score'] < baseline_m['combined_score']:
    winner = "🏆 BASELINE WINS"
    color = "red"
else:
    winner = "🤝 TIE"
    color = "gray"  
  

print(f"\n{winner}")  
  

detailed_question_analysis(baseline_results, gepa_results, test_questions)  
  

print("\n" + "="*80)
print("✓ COMPLETE RAG COMPARISON FINISHED")
print("=*80")
print("\nGenerated visualizations:")
print("  - rag_candidate_diversity.png")
print("  - rag_pareto_front_iter*.png")
print("  - rag_mutation_impact.png")
print("  - rag_merge_strategy.png")
print("  - rag_optimization_trajectory.png")
print("  - rag_comprehensive_comparison.png")

```

## Comprehensive RAG System Comparison: Baseline vs GEPA



=====

=

COMPARISON SUMMARY

=====

=

Retrieval Recall:

Baseline: 0.0000  
GEPA: 0.0000  
Improvement: +0.0000 (+0.00%)

Generation F1:

Baseline: 0.0987  
GEPA: 0.0751  
Improvement: -0.0236 (-23.92%)

Faithfulness:

Baseline: 0.8600  
GEPA: 0.9000  
Improvement: +0.0400 (+4.65%)

Combined Score:

Baseline: 0.2115  
GEPA: 0.2100  
Improvement: -0.0014 (-0.68%)

=====

=

Win/Loss/Tie: 2/3/0  
GEPA Win Rate: 40.0%

=====

=

=====

=

DETAILED QUESTION-BY-QUESTION ANALYSIS

=====

=

=====

=

QUESTION 1: What is the eye of the machine?

=====

=

Gold Answer: computer vision and pattern recognition systems

Baseline Answer: The "eye of the machine" refers to the capability of machines, especially those utilizing deep learning and neural networks, to process and understand visual information, akin to the human visual syst...

GEPA Answer: The "eye of the machine" likely refers to the ability of machines, particularly those using deep neural networks, to process and understand visual information, akin to the human eye. This concept is r...

Metric	Baseline	GEPA	Improvement
retrieval_recall	0.0000	0.0000	+0.0000 =
generation_f1	0.1404	0.0800	-0.0604 x
faithfulness	0.7000	0.9000	+0.2000 ✓
combined_score	0.1961	0.2120	+0.0159 ✓

🏆 GEPA WINS

=

QUESTION 2: How do machines learn?

=

Gold Answer: through algorithms that identify patterns in data

Baseline Answer: Machines learn through algorithms like the perceptron, which adjusts its predictions by evaluating the correctness of its outputs against given labels and modifying its weights and bias to minimize er...

GEPA Answer: Machines learn through algorithms that enable them to adjust their predictions based on the data they receive. A key model in this process is the perceptron, which modifies its weights and biases to m...

Metric	Baseline	GEPA	Improvement
retrieval_recall	0.0000	0.0000	+0.0000 =
generation_f1	0.2105	0.1449	-0.0656 x
faithfulness	0.9000	0.9000	+0.0000 =
combined_score	0.2642	0.2380	-0.0262 x

🏆 BASELINE WINS

=

QUESTION 3: What is the role of neural networks? what is the math behind it?

=

Gold Answer: neural networks process information through interconnected layers

Baseline Answer: Neural networks play a crucial role in machine learning, enabling computers to learn from data and improve their performance over time. The math behind neural networks involves concepts such as linear...

GEPA Answer: The role of neural networks is to learn to represent complex functions and make predictions or decisions based on input data. The math behind neural networks involves algorithms like backpropagation, ...

Metric	Baseline	GEPA	Improvement
retrieval_recall	0.0000	0.0000	+0.0000 =

generation_f1	0.0550	0.0714	+0.0164 ✓
faithfulness	0.9000	0.9000	+0.0000 =
combined_score	0.2020	0.2086	+0.0066 ✓

🏆 GEPA WINS

---



---

=  
QUESTION 4: What the author trying to convey from this text?  
=====

Gold Answer: overfitting, data quality, and generalization

Baseline Answer: The author is trying to convey the importance of critically evaluating information, considering alternative perspectives, and recognizing the complexity and nuance of various topics, including the lim...

GEPA Answer: The author is trying to convey that the development of artificial intelligence and related fields is a complex and multifaceted process, influenced by a variety of factors and perspectives, and that i...

Metric	Baseline	GEPA	Improvement
retrieval_recall	0.0000	0.0000	+0.0000 =
generation_f1	0.0465	0.0400	-0.0065 ✗
faithfulness	0.9000	0.9000	+0.0000 =
combined_score	0.1986	0.1960	-0.0026 ✗

🏆 BASELINE WINS

---



---

=  
QUESTION 5: AI evoloved and its history ?  
=====

Gold Answer: using metrics like accuracy, precision, recall, and F1 score

Baseline Answer: The evolution of AI has been marked by significant developments, from the early perceptrons to the current deep learning systems. The journey began with simple linear classifications and progressed to...

GEPA Answer: The evolution of AI has been a significant journey, from the early concept of perceptrons to the current complex deep learning systems. The development of AI has been fueled by advancements in mathema...

Metric	Baseline	GEPA	Improvement
retrieval_recall	0.0000	0.0000	+0.0000 =
generation_f1	0.0412	0.0392	-0.0020 ✗
faithfulness	0.9000	0.9000	+0.0000 =
combined_score	0.1965	0.1957	-0.0008 ✗

🏆 BASELINE WINS

```
=====
= ✓ COMPLETE RAG COMPARISON FINISHED =====
=
=
```

Generated visualizations:

- rag\_candidate\_diversity.png
- rag\_pareto\_front\_iter\*.png
- rag\_mutation\_impact.png
- rag\_merge\_strategy.png
- rag\_optimization\_trajectory.png
- rag\_comprehensive\_comparison.png