

# Task 1: Summarizer

```
In [2]: # Install required packages
```

```
!pip install dspy-ai PyPDF2 numpy matplotlib seaborn scikit-learn sentence-tr
```

Requirement already satisfied: dspy-ai in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (2.6.27)  
Requirement already satisfied: PyPDF2 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (3.0.1)  
Requirement already satisfied: numpy in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (2.2.6)  
Requirement already satisfied: matplotlib in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (3.10.7)  
Requirement already satisfied: seaborn in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (0.13.2)  
Requirement already satisfied: scikit-learn in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (1.7.2)  
Requirement already satisfied: sentence-transformers in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (5.1.2)  
Requirement already satisfied: pandas in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (2.2.3)  
Requirement already satisfied: nltk in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (3.9.2)  
Requirement already satisfied: dspy>=2.6.5 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy-ai) (2.6.5)  
Requirement already satisfied: contourpy>=1.0.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from matplotlib) (1.3.3)  
Requirement already satisfied: cycler>=0.10 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from matplotlib) (0.12.1)  
Requirement already satisfied: fonttools>=4.22.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from matplotlib) (4.60.1)  
Requirement already satisfied: kiwisolver>=1.3.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from matplotlib) (1.4.9)  
Requirement already satisfied: packaging>=20.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from matplotlib) (24.2)  
Requirement already satisfied: pillow>=8 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from matplotlib) (11.3.0)  
Requirement already satisfied: pyparsing>=3 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from matplotlib) (3.2.5)  
Requirement already satisfied: python-dateutil>=2.7 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from matplotlib) (2.9.0.post0)  
Requirement already satisfied: scipy>=1.8.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from scikit-learn) (1.16.3)  
Requirement already satisfied: joblib>=1.2.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from scikit-learn) (1.5.2)  
Requirement already satisfied: threadpoolctl>=3.1.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from scikit-learn) (3.6.0)  
Requirement already satisfied: transformers<5.0.0,>=4.41.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from sentence-transformers) (4.57.3)  
Requirement already satisfied: tqdm in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from sentence-transformers) (4.67.1)  
Requirement already satisfied: torch>=1.11.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from sentence-transformers) (2.9.1)  
Requirement already satisfied: huggingface-hub>=0.20.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from sentence-transformers) (0.35.3)  
Requirement already satisfied: typing\_extensions>=4.5.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from sentence-transformers)

(4.15.0)

Requirement already satisfied: filelock in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from transformers<5.0.0,>=4.41.0->sentence-transformers) (3.20.0)

Requirement already satisfied: pyyaml>=5.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from transformers<5.0.0,>=4.41.0->sentence-transformers) (6.0.3)

Requirement already satisfied: regex!=2019.12.17 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from transformers<5.0.0,>=4.41.0->sentence-transformers) (2025.9.18)

Requirement already satisfied: requests in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from transformers<5.0.0,>=4.41.0->sentence-transformers) (2.32.5)

Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from transformers<5.0.0,>=4.41.0->sentence-transformers) (0.22.1)

Requirement already satisfied: safetensors>=0.4.3 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from transformers<5.0.0,>=4.41.0->sentence-transformers) (0.7.0)

Requirement already satisfied: fsspec>=2023.5.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from huggingface-hub>=0.20.0->sentence-transformers) (2025.9.0)

Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from huggingface-hub>=0.20.0->sentence-transformers) (1.1.10)

Requirement already satisfied: pytz>=2020.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from pandas) (2025.2)

Requirement already satisfied: tzdata>=2022.7 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from pandas) (2025.2)

Requirement already satisfied: click in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from nltk) (8.3.0)

Requirement already satisfied: anyio in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (4.11.0)

Requirement already satisfied: asyncc==0.0.8 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (0.0.8)

Requirement already satisfied: backoff in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (2.2.1)

Requirement already satisfied: cachetools in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (6.2.0)

Requirement already satisfied: cloudpickle in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (3.1.2)

Requirement already satisfied: datasets in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (4.4.1)

Requirement already satisfied: diskcache in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (5.6.3)

Requirement already satisfied: httpx in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (0.28.1)

Requirement already satisfied: jinja2 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (3.1.6)

Requirement already satisfied: json-repair in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (0.54.2)

Requirement already satisfied: litellm<2.0.0,>=1.59.8 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.80.7)

```
Requirement already satisfied: magicattr~=0.1.6 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (0.1.6)
Requirement already satisfied: openai in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (2.9.0)
Requirement already satisfied: optuna in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (4.6.0)
Requirement already satisfied: pydantic~=2.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (2.12.0)
Requirement already satisfied: tenacity>=8.2.3 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (9.1.2)
Requirement already satisfied: ujson in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from dspy>=2.6.5->dspy-ai) (5.11.0)
Requirement already satisfied: idna>=2.8 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from anyio->dspy>=2.6.5->dspy-ai) (3.10)
Requirement already satisfied: sniffio>=1.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from anyio->dspy>=2.6.5->dspy-ai) (1.3.1)
Requirement already satisfied: aiohttp>=3.10 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (3.13.0)
Requirement already satisfied: fastuuid>=0.13.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.14.0)
Requirement already satisfied: grpcio<1.68.0,>=1.62.3 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.67.1)
Requirement already satisfied: importlib-metadata>=6.8.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (8.7.0)
Requirement already satisfied: jsonschema<5.0.0,>=4.22.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (4.25.1)
Requirement already satisfied: python-dotenv>=0.2.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.1.1)
Requirement already satisfied: tiktoken>=0.7.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.12.0)
Requirement already satisfied: MarkupSafe>=2.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from jinja2->dspy>=2.6.5->dspy-ai) (3.0.3)
Requirement already satisfied: attrs>=22.2.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from jsonschema<5.0.0,>=4.22.0->litellm<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (25.4.0)
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from jsonschema<5.0.0,>=4.22.0->litellm<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (2025.9.1)
Requirement already satisfied: referencing>=0.28.4 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from jsonschema<5.0.0,>=4.22.0->litellm<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.37.0)
Requirement already satisfied: rpds-py>=0.7.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from jsonschema<5.0.0,>=4.22.0->litellm<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.30.0)
```

Requirement already satisfied: PyJWT<3.0.0,>=2.10.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (2.10.1)

Requirement already satisfied: apscheduler<4.0.0,>=3.10.4 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (3.11.1)

Requirement already satisfied: azure-identity<2.0.0,>=1.15.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.25.1)

Requirement already satisfied: azure-storage-blob<13.0.0,>=12.25.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (12.27.1)

Requirement already satisfied: boto3==1.36.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.36.0)

Requirement already satisfied: cryptography in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (46.0.2)

Requirement already satisfied: fastapi>=0.120.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.124.0)

Requirement already satisfied: fastapi-sso<0.17.0,>=0.16.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.16.0)

Requirement already satisfied: gunicorn<24.0.0,>=23.0.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (23.0.0)

Requirement already satisfied: litellm-enterprise==0.1.22 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.1.22)

Requirement already satisfied: litellm-proxy-extras==0.4.9 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.4.9)

Requirement already satisfied: mcp<2.0.0,>=1.21.2 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.23.1)

Requirement already satisfied: orjson<4.0.0,>=3.9.7 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (3.11.3)

Requirement already satisfied: polars<2.0.0,>=1.31.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.35.2)

Requirement already satisfied: pynacl<2.0.0,>=1.5.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.6.1)

Requirement already satisfied: python-multipart<0.0.19,>=0.0.18 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.0.18)

Requirement already satisfied: rich==13.7.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (13.7.1)

Requirement already satisfied: rq in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (2.6.1)

```
Requirement already satisfied: soundfile<0.13.0,>=0.12.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.12.1)
Requirement already satisfied: uvicorn<0.32.0,>=0.31.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.31.1)
Requirement already satisfied: uvloop<0.22.0,>=0.21.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.21.0)
Requirement already satisfied: websockets<16.0.0,>=15.0.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (15.0.1)
Requirement already satisfied: botocore<1.37.0,>=1.36.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from boto3==1.36.0->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.36.26)
Requirement already satisfied: jmespath<2.0.0,>=0.7.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from boto3==1.36.0->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.0.1)
Requirement already satisfied: s3transfer<0.12.0,>=0.11.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from boto3==1.36.0->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.11.3)
Requirement already satisfied: markdown-it-py>=2.2.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from rich==13.7.1->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (4.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from rich==13.7.1->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (2.19.2)
Requirement already satisfied: tzlocal>=3.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from apscheduler<4.0.0,>=3.10.4->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (5.3.1)
Requirement already satisfied: azure-core>=1.31.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from azure-identity<2.0.0,>=1.15.0->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.36.0)
Requirement already satisfied: msal>=1.30.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from azure-identity<2.0.0,>=1.15.0->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.34.0)
Requirement already satisfied: msal-extensions>=1.2.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from azure-identity<2.0.0,>=1.15.0->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.3.1)
Requirement already satisfied: isodate>=0.6.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from azure-storage-blob<13.0.0,>=12.25.1->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.7.2)
Requirement already satisfied: urllib3!=2.2.0,<3,>=1.25.4 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from botocore<1.37.0,>=1.36.0->boto3==1.36.0->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (2.5.0)
Requirement already satisfied: oauthlib>=3.1.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from fastapi-sso<0.17.0,>=0.16.0->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (3.3.1)
Requirement already satisfied: httpx-sse>=0.4 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from mcp<2.0.0,>=1.21.2->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.4.2)
Requirement already satisfied: pydantic-settings>=2.5.2 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from mcp<2.0.0,>=1.21.2->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (2.5.2)
```

```
lm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (2.11.0)
Requirement already satisfied: sse-starlette>=1.6.1 in /home/kelvin/miniconda3/
envs/genai-dev/lib/python3.14/site-packages (from mcp<2.0.0,>=1.21.2->litellm[p
roxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (3.0.3)
Requirement already satisfied: starlette>=0.27 in /home/kelvin/miniconda3/envs/
genai-dev/lib/python3.14/site-packages (from mcp<2.0.0,>=1.21.2->litellm[prox
y]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.48.0)
Requirement already satisfied: typing-inspection>=0.4.1 in /home/kelvin/minicon
da3/envs/genai-dev/lib/python3.14/site-packages (from mcp<2.0.0,>=1.21.2->litel
lm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.4.2)
Requirement already satisfied: polars-runtime-32==1.35.2 in /home/kelvin/minico
nda3/envs/genai-dev/lib/python3.14/site-packages (from polars<2.0.0,>=1.31.0->l
itellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.35.2)
Requirement already satisfied: annotated-types>=0.6.0 in /home/kelvin/miniconda
3/envs/genai-dev/lib/python3.14/site-packages (from pydantic~2.0->dsp
y>=2.6.5->dspy-ai) (0.7.0)
Requirement already satisfied: pydantic-core==2.41.1 in /home/kelvin/miniconda
3/envs/genai-dev/lib/python3.14/site-packages (from pydantic~2.0->dsp
y>=2.6.5->dspy-ai) (2.41.1)
Requirement already satisfied: cffi>=2.0.0 in /home/kelvin/miniconda3/envs/gena
i-dev/lib/python3.14/site-packages (from pynacl<2.0.0,>=1.5.0->litellm[prox
y]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (2.0.0)
Requirement already satisfied: six>=1.5 in /home/kelvin/miniconda3/envs/genai-d
ev/lib/python3.14/site-packages (from python-dateutil>=2.7->matplotlib) (1.1
7.0)
Requirement already satisfied: h11>=0.8 in /home/kelvin/miniconda3/envs/genai-d
ev/lib/python3.14/site-packages (from uvicorn<0.32.0,>=0.31.1->litellm[prox
y]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.16.0)
Requirement already satisfied: aiohappyeyeballs>=2.5.0 in /home/kelvin/minicond
a3/envs/genai-dev/lib/python3.14/site-packages (from aiohttp>=3.10->litell
m<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (2.6.1)
Requirement already satisfied: aiosignal>=1.4.0 in /home/kelvin/miniconda3/env
s/genai-dev/lib/python3.14/site-packages (from aiohttp>=3.10->litell
m<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.4.0)
Requirement already satisfied: frozenlist>=1.1.1 in /home/kelvin/miniconda3/env
s/genai-dev/lib/python3.14/site-packages (from aiohttp>=3.10->litell
m<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.8.0)
Requirement already satisfied: multidict<7.0,>=4.5 in /home/kelvin/miniconda3/e
nvs/genai-dev/lib/python3.14/site-packages (from aiohttp>=3.10->litell
m<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (6.7.0)
Requirement already satisfied: propcache>=0.2.0 in /home/kelvin/miniconda3/env
s/genai-dev/lib/python3.14/site-packages (from aiohttp>=3.10->litell
m<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.4.1)
Requirement already satisfied: yarl<2.0,>=1.17.0 in /home/kelvin/miniconda3/env
s/genai-dev/lib/python3.14/site-packages (from aiohttp>=3.10->litell
m<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (1.22.0)
Requirement already satisfied: pycparser in /home/kelvin/miniconda3/envs/genai-
dev/lib/python3.14/site-packages (from cffi>=2.0.0->pynacl<2.0.0,>=1.5.0->litel
lm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (2.23)
Requirement already satisfied: annotated-doc>=0.0.2 in /home/kelvin/miniconda3/
envs/genai-dev/lib/python3.14/site-packages (from fastapi>=0.120.1->litellm[pro
xy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.0.4)
Requirement already satisfied: certifi in /home/kelvin/miniconda3/envs/genai-de
v/lib/python3.14/site-packages (from httpx->dspy>=2.6.5->dspy-ai) (2025.10.5)
```

Requirement already satisfied: httpcore==1.\* in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from httpx->dspy>=2.6.5->dspy-ai) (1.0.9)

Requirement already satisfied: zipp>=3.20 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from importlib-metadata>=6.8.0->litellm<2.0.0,>=1.59.8->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (3.23.0)

Requirement already satisfied: mdurl~0.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from markdown-it-py>=2.2.0->rich==13.7.1->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (0.1.2)

Requirement already satisfied: charset\_normalizer<4,>=2 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from requests->transformers<5.0.0,>=4.41.0->sentence-transformers) (3.4.3)

Requirement already satisfied: distro<2,>=1.7.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from openai->dspy>=2.6.5->dspy-ai) (1.9.0)

Requirement already satisfied: jiter<1,>=0.10.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from openai->dspy>=2.6.5->dspy-ai) (0.11.0)

Requirement already satisfied: email-validator>=2.0.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from pydantic[email]>=1.8.0->fastapi-sso<0.17.0,>=0.16.0->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (2.3.0)

Requirement already satisfied: dnspython>=2.0.0 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from email-validator>=2.0.0->pydantic[email]>=1.8.0->fastapi-sso<0.17.0,>=0.16.0->litellm[proxy]<2.0.0,>=1.59.8->dspy>=2.6.5->dspy-ai) (2.8.0)

Requirement already satisfied: setuptools in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentence-transformers) (80.9.0)

Requirement already satisfied: sympy>=1.13.3 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentence-transformers) (1.14.0)

Requirement already satisfied: networkx>=2.5.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentence-transformers) (3.5)

Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.8.93 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentence-transformers) (12.8.93)

Requirement already satisfied: nvidia-cuda-runtime-cu12==12.8.90 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentence-transformers) (12.8.90)

Requirement already satisfied: nvidia-cuda-cupti-cu12==12.8.90 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentence-transformers) (12.8.90)

Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentence-transformers) (9.10.2.21)

Requirement already satisfied: nvidia-cublas-cu12==12.8.4.1 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentence-transformers) (12.8.4.1)

Requirement already satisfied: nvidia-cufft-cu12==11.3.3.83 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentence-transformers) (11.3.3.83)

Requirement already satisfied: nvidia-curand-cu12==10.3.9.90 in /home/kelvin/miniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sente

```
nce-transformers) (10.3.9.90)
Requirement already satisfied: nvidia-cusolver-cu12==11.7.3.90 in /home/kelvin/
miniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sen
tence-transformers) (11.7.3.90)
Requirement already satisfied: nvidia-cusparse-cu12==12.5.8.93 in /home/kelvin/
miniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sen
tence-transformers) (12.5.8.93)
Requirement already satisfied: nvidia-cusparseelt-cu12==0.7.1 in /home/kelvin/mi
niconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sente
nce-transformers) (0.7.1)
Requirement already satisfied: nvidia-nccl-cu12==2.27.5 in /home/kelvin/minicon
da3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentence-t
ravers) (2.27.5)
Requirement already satisfied: nvidia-nvshmem-cu12==3.3.20 in /home/kelvin/minicon
da3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentencec
e-transformers) (3.3.20)
Requirement already satisfied: nvidia-nvtx-cu12==12.8.90 in /home/kelvin/minico
nda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentence-
transformers) (12.8.90)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.8.93 in /home/kelvin/m
iniconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sent
ence-transformers) (12.8.93)
Requirement already satisfied: nvidia-cufile-cu12==1.13.1.3 in /home/kelvin/min
iconda3/envs/genai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sente
nce-transformers) (1.13.1.3)
Requirement already satisfied: triton==3.5.1 in /home/kelvin/miniconda3/envs/ge
nai-dev/lib/python3.14/site-packages (from torch>=1.11.0->sentence-transformer
s) (3.5.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /home/kelvin/miniconda3/en
vs/genai-dev/lib/python3.14/site-packages (from sympy>=1.13.3->torch>=1.11.0->s
entence-transformers) (1.3.0)
Requirement already satisfied: pyarrow>=21.0.0 in /home/kelvin/miniconda3/envs/
genai-dev/lib/python3.14/site-packages (from datasets->dspy>=2.6.5->dspy-ai) (2
2.0.0)
Requirement already satisfied: dill<0.4.1,>=0.3.0 in /home/kelvin/miniconda3/en
vs/genai-dev/lib/python3.14/site-packages (from datasets->dspy>=2.6.5->dspy-ai)
(0.4.0)
Requirement already satisfied: xxhash in /home/kelvin/miniconda3/envs/genai-de
v/lib/python3.14/site-packages (from datasets->dspy>=2.6.5->dspy-ai) (3.6.0)
Requirement already satisfied: multiprocessing<0.70.19 in /home/kelvin/miniconda3/
envs/genai-dev/lib/python3.14/site-packages (from datasets->dspy>=2.6.5->dspy-a
i) (0.70.18)
Requirement already satisfied: alembic>=1.5.0 in /home/kelvin/miniconda3/envs/g
enai-dev/lib/python3.14/site-packages (from optuna->dspy>=2.6.5->dspy-ai) (1.1
7.2)
Requirement already satisfied: colorlog in /home/kelvin/miniconda3/envs/genai-d
ev/lib/python3.14/site-packages (from optuna->dspy>=2.6.5->dspy-ai) (6.10.1)
Requirement already satisfied: sqlalchemy>=1.4.2 in /home/kelvin/miniconda3/en
vs/genai-dev/lib/python3.14/site-packages (from optuna->dspy>=2.6.5->dspy-ai)
(2.0.43)
Requirement already satisfied: Mako in /home/kelvin/miniconda3/envs/genai-dev/l
ib/python3.14/site-packages (from alembic>=1.5.0->optuna->dspy>=2.6.5->dspy-ai)
(1.3.10)
Requirement already satisfied: croniter in /home/kelvin/miniconda3/envs/genai-d
```

```
ev/lib/python3.14/site-packages (from rq->litellm[proxy]<2.0.0,>=1.59.8->dsp  
y>=2.6.5->dspy-ai) (6.0.0)  
Requirement already satisfied: redis!=6,>=3.5 in /home/kelvin/miniconda3/envs/g  
enai-dev/lib/python3.14/site-packages (from rq->litellm[proxy]<2.0.0,>=1.5  
9.8->dspy>=2.6.5->dspy-ai) (7.1.0)
```

```
In [3]: # --- Standard Library ---  
import getpass  
import re  
import random  
from typing import List, Dict, Tuple  
from dataclasses import dataclass  
from collections import defaultdict  
  
# --- Third-Party Libraries ---  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
import textwrap  
  
# Machine Learning / NLP  
import dspy  
import PyPDF2  
from sklearn.metrics.pairwise import cosine_similarity  
from sentence_transformers import SentenceTransformer  
  
# --- NLTK (Tokenization, Stopwords, Lemmatization) ---  
import nltk  
nltk.download('punkt', quiet=True)  
nltk.download('stopwords', quiet=True)  
nltk.download('wordnet', quiet=True)  
nltk.download('omw-1.4', quiet=True)  
nltk.download('punkt_tab', quiet=True)  
  
from nltk.corpus import stopwords  
from nltk.stem import WordNetLemmatizer  
  
STOPWORDS = set(stopwords.words('english'))  
lemmatizer = WordNetLemmatizer()  
  
# --- Random Seed Setup ---  
RANDOM_STATE = 42  
np.random.seed(RANDOM_STATE)  
random.seed(RANDOM_STATE)  
  
# --- Visualization Defaults ---  
sns.set_style("whitegrid")  
plt.rcParams['figure.figsize'] = (12, 6)
```

LLM INITIALIZATION (If Using Grok)

```
In [4]: %pip install getpass langchain_groq langchain_core
```

```
ERROR: Could not find a version that satisfies the requirement getpass (from ve
rsions: none)
ERROR: No matching distribution found for getpass
Note: you may need to restart the kernel to use updated packages.
```

```
In [5]: api_key = getpass.getpass("Enter your GROQ API key: ")

from langchain_groq import ChatGroq
from langchain_core.messages import HumanMessage, SystemMessage

class GroqlM(dspy.LM):
    def __init__(self, model="llama-3.3-70b-versatile", api_key=None, **kwargs):
        super().__init__(model=model, **kwargs)
        self.model_name = model
        self.client = ChatGroq(model=model, api_key=api_key)

    def __call__(self, prompt=None, messages=None, **kwargs):
        if messages:
            lc_messages = []
            for msg in messages:
                if msg.get("role") == "system":
                    lc_messages.append(SystemMessage(content=msg["content"]))
                elif msg.get("role") == "user":
                    lc_messages.append(HumanMessage(content=msg["content"]))
            response = self.client.invoke(lc_messages)
        else:
            response = self.client.invoke([HumanMessage(content=prompt)])
        return [response.content]

# Initialize LM (use your actual API key)
lm = GroqlM(api_key=api_key)
dspy.configure(lm=lm)
```

## PDF EXTRACTION UTILITY

```
In [6]: def extract_pdf_text(path):
    """Extract text from PDF file"""
    text = ""
    with open(path, "rb") as f:
        reader = PyPDF2.PdfReader(f)
        for page in reader.pages:
            text += page.extract_text() + "\n"
    return text
```

## BASELINE DSPY SUMMARIZER (Original)

```
In [7]: class SummarizePDF(dspy.Signature):
    """Summarize a long document extracted from a PDF with quality metadata."""
    document: str = dspy.InputField()
    summary: str = dspy.OutputField(
        desc="Concise, high-level summary of the document."
    )
```

```

key_points: list[str] = dspy.OutputField(
    desc="A bullet list of the most important points."
)
confidence: float = dspy.OutputField(
    desc="Model-estimated confidence score between 0 and 1."
)
length_tokens: int = dspy.OutputField(
    desc="Approximate token length of the summary."
)
coverage_score: float = dspy.OutputField(
    desc="Model's self-rating of how well the summary covers the document"
)

class BaselinePDFSummarizer(dspy.Module):
    def __init__(self):
        super().__init__()
        self.summarize = dspy.ChainOfThought(SummarizePDF)

    def forward(self, document: str):
        out = self.summarize(document=document)
        return dspy.Prediction(
            summary=out.summary,
            key_points=out.key_points,
            confidence=out.confidence,
            length_tokens=out.length_tokens,
            coverage_score=out.coverage_score
        )

```

## GEPA COMPONENTS - PROMPT CANDIDATE STRUCTURE

```

In [9]: from dataclasses import dataclass, field

@dataclass
class PromptCandidate:
    """Represents a single prompt candidate in the GEPA system"""
    prompt_template: str
    instruction_prefix: str
    output_format: str
    reasoning_style: str
    score: float = 0.0
    performance_metrics: Dict[str, float] = field(default_factory=dict)

    def to_dict(self):
        return {
            'prompt_template': self.prompt_template,
            'instruction_prefix': self.instruction_prefix,
            'output_format': self.output_format,
            'reasoning_style': self.reasoning_style,
            'score': self.score,
            'performance_metrics': self.performance_metrics
        }

```

## GEPA - CANDIDATE POOL INITIALIZATION

```
In [10]: class CandidatePoolInitializer:
    """Initialize diverse pool of prompt candidates"""

    def __init__(self, pool_size: int = 10):
        self.pool_size = pool_size
        self.embedding_model = SentenceTransformer('all-MiniLM-L6-v2')

    def generate_initial_pool(self) -> List[PromptCandidate]:
        """Generate diverse initial candidate pool"""

        # Different instruction prefixes
        prefixes = [
            "Analyze the following document carefully and",
            "Read through this document thoroughly, then",
            "Given the document below, your task is to",
            "Extract key information from this document and",
            "Comprehensively review this document to",
        ]

        # Different output formats
        formats = [
            "structured summary with clear sections",
            "hierarchical breakdown of main concepts",
            "narrative summary with logical flow",
            "analytical digest with critical insights",
        ]

        # Different reasoning styles
        reasoning = [
            "step-by-step analytical approach",
            "holistic synthesis of key themes",
            "critical evaluation of main arguments",
            "systematic extraction of core concepts",
        ]

        # Different prompt templates
        templates = [
            "{prefix} provide a {format} using {reasoning}.",
            "{prefix} generate a {format}. Apply {reasoning} to ensure complet",
            "Using {reasoning}, {prefix} create a {format}.",
            "{prefix} deliver a {format}. Ensure {reasoning} is applied throug",
        ]

        candidates = []
        np.random.seed(42)

        for i in range(self.pool_size):
            candidate = PromptCandidate(
                prompt_template=np.random.choice(templates),
                instruction_prefix=np.random.choice(prefixes),
                output_format=np.random.choice(formats),
                reasoning_style=np.random.choice(reasoning)
            )
            candidates.append(candidate)
        return candidates
```

```

        candidates.append(candidate)

    return candidates

def visualize_pool_diversity(self, candidates: List[PromptCandidate]):
    """Visualize diversity of candidate pool"""

    # Create embeddings for each candidate
    texts = [f"{c.instruction_prefix} {c.output_format} {c.reasoning_style}"
             for c in candidates]
    embeddings = self.embedding_model.encode(texts)

    # Compute pairwise similarities
    similarities = cosine_similarity(embeddings)

    # Plot heatmap
    plt.figure(figsize=(10, 8))
    sns.heatmap(similarities, annot=True, fmt='.2f', cmap='coolwarm',
                xticklabels=[f"C{i+1}" for i in range(len(candidates))],
                yticklabels=[f"C{i+1}" for i in range(len(candidates))])
    plt.title("Candidate Pool Diversity (Cosine Similarity Matrix)")
    plt.tight_layout()
    plt.savefig('candidate_pool_diversity.png', dpi=300, bbox_inches='tight')
    plt.show()

    print(f"\n✓ Generated {len(candidates)} diverse candidates")
    print(f"  Average pairwise similarity: {similarities[np.triu_indices_f

```

## GEPA - PARETO-BASED CANDIDATE FILTERING

```

In [11]: class ParetoFilter:
    """Filter candidates using Pareto dominance"""

    def __init__(self, objectives: List[str]):
        self.objectives = objectives

    def dominates(self, cand_a: PromptCandidate, cand_b: PromptCandidate) -> bool:
        """Check if candidate A Pareto-dominates candidate B"""
        metrics_a = cand_a.performance_metrics
        metrics_b = cand_b.performance_metrics

        better_in_any = False
        worse_in_any = False

        for obj in self.objectives:
            val_a = metrics_a.get(obj, 0)
            val_b = metrics_b.get(obj, 0)

            if val_a > val_b:
                better_in_any = True
            elif val_a < val_b:
                worse_in_any = True

        return better_in_any and not worse_in_any

```

```

        return better_in_any and not worse_in_any

    def get_pareto_front(self, candidates: List[PromptCandidate]) -> List[Prom
        """Extract Pareto-optimal candidates"""
        pareto_front = []

        for cand in candidates:
            is_dominated = False
            for other in candidates:
                if cand != other and self.dominates(other, cand):
                    is_dominated = True
                    break

            if not is_dominated:
                pareto_front.append(cand)

        return pareto_front

    def visualize_pareto_front(self, candidates: List[PromptCandidate],
                               pareto_front: List[PromptCandidate]):
        """Visualize Pareto front in 2D objective space"""

        if len(self.objectives) < 2:
            print("Need at least 2 objectives for visualization")
            return

        obj1, obj2 = self.objectives[0], self.objectives[1]

        # Extract metrics
        all_x = [c.performance_metrics.get(obj1, 0) for c in candidates]
        all_y = [c.performance_metrics.get(obj2, 0) for c in candidates]

        pareto_x = [c.performance_metrics.get(obj1, 0) for c in pareto_front]
        pareto_y = [c.performance_metrics.get(obj2, 0) for c in pareto_front]

        # Plot
        plt.figure(figsize=(10, 6))
        plt.scatter(all_x, all_y, c='lightblue', s=100, alpha=0.6,
                   label='All Candidates', edgecolors='black')
        plt.scatter(pareto_x, pareto_y, c='red', s=150, alpha=0.8,
                   label='Pareto Front', marker='*', edgecolors='darkred', linewidth=2)

        plt.xlabel(obj1.replace('_', ' ').title(), fontsize=12)
        plt.ylabel(obj2.replace('_', ' ').title(), fontsize=12)
        plt.title('Pareto Front Selection', fontsize=14, fontweight='bold')
        plt.legend(fontsize=10)
        plt.grid(True, alpha=0.3)
        plt.tight_layout()
        plt.savefig('pareto_front.png', dpi=300, bbox_inches='tight')
        plt.show()

    print(f"\n✓ Pareto front: {len(pareto_front)}/{len(candidates)} candidat

```

## GEPA - REFLECTIVE PROMPT MUTATION

```
In [12]: class ReflectiveMutator:
    """Mutate prompts based on performance reflection"""

    def __init__(self, lm):
        self.lm = lm
        self.mutation_history = []

    def analyze_performance(self, candidate: PromptCandidate) -> Dict[str, str]:
        """Analyze what worked and what didn't"""
        metrics = candidate.performance_metrics

        analysis = {
            'strengths': [],
            'weaknesses': [],
            'suggestions': []
        }

        # Analyze each metric
        if metrics.get('relevance', 0) > 0.7:
            analysis['strengths'].append("High relevance to source material")
        else:
            analysis['weaknesses'].append("Low relevance to source material")
            analysis['suggestions'].append("Enhance instruction to focus on key points")

        if metrics.get('coherence', 0) > 0.7:
            analysis['strengths'].append("Good logical flow")
        else:
            analysis['weaknesses'].append("Poor coherence")
            analysis['suggestions'].append("Add explicit structure requirements")

        if metrics.get('conciseness', 0) < 0.5:
            analysis['weaknesses'].append("Too verbose")
            analysis['suggestions'].append("Add length constraints")

    return analysis # type: ignore

    def mutate_candidate(self, candidate: PromptCandidate) -> PromptCandidate:
        """Generate mutated version based on performance analysis"""
        analysis = self.analyze_performance(candidate)

        # Simple rule-based mutations (in practice, could use LM for this)
        new_prefix = candidate.instruction_prefix
        new_format = candidate.output_format
        new_reasoning = candidate.reasoning_style

        # Apply mutations based on weaknesses
        if "Low relevance" in str(analysis['weaknesses']):
            new_prefix = "Carefully extract the most relevant information from"

        if "Poor coherence" in str(analysis['weaknesses']):
            new_format = "well-structured summary with clear logical progression"
```

```

        if "Too verbose" in str(analysis['weaknesses']):
            new_reasoning = "concise yet comprehensive analytical approach"

        mutated = PromptCandidate(
            prompt_template=candidate.prompt_template,
            instruction_prefix=new_prefix,
            output_format=new_format,
            reasoning_style=new_reasoning
        )

        self.mutation_history.append({
            'original': candidate.to_dict(),
            'mutated': mutated.to_dict(),
            'reason': analysis
        })

    return mutated

def visualize_mutation_history(self):
    """Visualize mutation improvements"""
    if not self.mutation_history:
        print("No mutations yet")
        return

    generations = list(range(len(self.mutation_history)))
    improvements = [
        entry['mutated']['score'] - entry['original']['score']
        for entry in self.mutation_history
    ]

    plt.figure(figsize=(10, 5))
    plt.bar(generations, improvements, color=['green' if x > 0 else 'red'])
    plt.axhline(y=0, color='black', linestyle='--', linewidth=1)
    plt.xlabel('Mutation Generation', fontsize=12)
    plt.ylabel('Score Improvement', fontsize=12)
    plt.title('Reflective Mutation Performance', fontsize=14, fontweight='bold')
    plt.tight_layout()
    plt.savefig('mutation_history.png', dpi=300, bbox_inches='tight')
    plt.show()

```

## GEPA - SYSTEM-AWARE MERGE

```

In [13]: class SystemAwareMerger:
    """Intelligently merge best candidates"""

    def __init__(self):
        self.merge_history = []

    def merge_candidates(self, candidates: List[PromptCandidate],
                         top_k: int = 3) -> PromptCandidate:
        """Merge top-k candidates into single optimized prompt"""

```

```

# Sort by score
sorted_cands = sorted(candidates, key=lambda x: x.score, reverse=True)

# Extract best components
best_prefix = self._select_best_component(
    [c.instruction_prefix for c in sorted_cands],
    [c.score for c in sorted_cands]
)

best_format = self._select_best_component(
    [c.output_format for c in sorted_cands],
    [c.score for c in sorted_cands]
)

best_reasoning = self._select_best_component(
    [c.reasoning_style for c in sorted_cands],
    [c.score for c in sorted_cands]
)

# Use template from best candidate
best_template = sorted_cands[0].prompt_template

merged = PromptCandidate(
    prompt_template=best_template,
    instruction_prefix=best_prefix,
    output_format=best_format,
    reasoning_style=best_reasoning,
    score=sorted_cands[0].score # Inherit best score
)

self.merge_history.append({
    'sources': [c.to_dict() for c in sorted_cands],
    'merged': merged.to_dict()
})

return merged

def _select_best_component(self, components: List[str], scores: List[float]):
    """Select best component weighted by scores"""
    # Simple: return component from best candidate
    best_idx = np.argmax(scores)
    return components[best_idx]

def visualize_merge_process(self):
    """Visualize how merging improves performance"""
    if not self.merge_history:
        print("No merges yet")
        return

    merge_rounds = []
    avg_source_scores = []
    merged_scores = []


```

```

for i, entry in enumerate(self.merge_history):
    merge_rounds.append(f"Merge {i+1}")
    avg_source = np.mean([s['score'] for s in entry['sources']])
    avg_source_scores.append(avg_source)
    merged_scores.append(entry['merged']['score'])

x = np.arange(len(merge_rounds))
width = 0.35

fig, ax = plt.subplots(figsize=(10, 6))
ax.bar(x - width/2, avg_source_scores, width, label='Avg Source Score')
ax.bar(x + width/2, merged_scores, width, label='Merged Score', color='red')

ax.set_xlabel('Merge Round', fontsize=12)
ax.set_ylabel('Score', fontsize=12)
ax.set_title('System-Aware Merge Performance', fontsize=14, fontweight='bold')
ax.set_xticks(x)
ax.set_xticklabels(merge_rounds)
ax.legend()
plt.tight_layout()
plt.savefig('merge_performance.png', dpi=300, bbox_inches='tight')
plt.show()

```

## GEPA SUMMARIZER (Main Class)

```

In [14]: class GEPAsummarizer(dspy.Module):
    """GEPA-optimized PDF Summarizer"""

    def __init__(self, lm, n_iterations: int = 3):
        super().__init__()
        self.lm = lm
        self.n_iterations = n_iterations

        # GEPA components
        self.pool_initializer = CandidatePoolInitializer(pool_size=10)
        self.pareto_filter = ParetoFilter(objectives=['relevance', 'coherence'])
        self.mutator = ReflectiveMutator(lm)
        self.merger = SystemAwareMerger()

        # Best candidate
        self.optimized_prompt = None
        self.optimization_history = []

    def optimize_prompts(self, sample_documents: List[str]):
        """Run GEPA optimization loop"""

        print("=*70")
        print("GEPA PROMPT OPTIMIZATION")
        print("=*70")

        # Step 1: Initialize candidate pool
        print("\n[Step 1] Initializing Candidate Pool...")
        candidates = self.pool_initializer.generate_initial_pool()

```

```

        self.pool_initializer.visualize_pool_diversity(candidates)

    # Step 2: Evaluate candidates on sample documents
    print("\n[Step 2] Evaluating Initial Candidates...")
    for candidate in candidates:
        metrics = self._evaluate_candidate(candidate, sample_documents[0])
        candidate.performance_metrics = metrics
        candidate.score = np.mean(list(metrics.values())) # type: ignore

    # Optimization loop
    for iteration in range(self.n_iterations):
        print(f"\n{'='*70}")
        print(f"ITERATION {iteration + 1}/{self.n_iterations}")
        print(f"{'='*70}")

    # Step 3: Pareto filtering
    print(f"\n[Step 3.{iteration+1}] Applying Pareto Filtering...")
    pareto_front = self.pareto_filter.get_pareto_front(candidates)
    self.pareto_filter.visualize_pareto_front(candidates, pareto_front)

    # Step 4: Reflective mutation
    print(f"\n[Step 4.{iteration+1}] Applying Reflective Mutations...")
    mutated = []
    for candidate in pareto_front[:5]: # Mutate top 5
        mutated_candidate = self.mutator.mutate_candidate(candidate)
        metrics = self._evaluate_candidate(mutated_candidate, sample_documents[0])
        mutated_candidate.performance_metrics = metrics
        mutated_candidate.score = np.mean(list(metrics.values())) # type: ignore
        mutated.append(mutated_candidate)

    self.mutator.visualize_mutation_history()

    # Step 5: System-aware merge
    print(f"\n[Step 5.{iteration+1}] Merging Best Candidates...")
    merged = self.merger.merge_candidates(pareto_front + mutated, top_k=5)

    # Re-evaluate merged
    merged_metrics = self._evaluate_candidate(merged, sample_documents[0])
    merged.performance_metrics = merged_metrics
    merged.score = np.mean(list(merged_metrics.values())) # type: ignore

    self.merger.visualize_merge_process()

    # Update candidate pool
    candidates = pareto_front + mutated + [merged]

    # Track history
    self.optimization_history.append({
        'iteration': iteration + 1,
        'best_score': max(c.score for c in candidates),
        'avg_score': np.mean([c.score for c in candidates]),
        'pareto_size': len(pareto_front)
    })

```

```

# Select final best candidate
self.optimized_prompt = max(candidates, key=lambda x: x.score)
print(f"\n Optimization Complete!")
print(f" Final Best Score: {self.optimized_prompt.score:.3f}")

# Visualize optimization trajectory
self._visualize_optimization_trajectory()

def _evaluate_candidate(self, candidate: PromptCandidate, document: str) ->
    """Evaluate a candidate on a document"""

    # Build full prompt
    full_prompt = candidate.prompt_template.format(
        prefix=candidate.instruction_prefix,
        format=candidate.output_format,
        reasoning=candidate.reasoning_style
    )

    full_prompt += f"\n\nDocument:\n{document[:2000]}" # Truncate for speed

    # Simulate metrics (in practice, use actual LM + evaluation)
    # Here we use heuristics for demonstration
    metrics = {
        'relevance': np.random.uniform(0.6, 0.9),
        'coherence': np.random.uniform(0.6, 0.9),
        'conciseness': np.random.uniform(0.5, 0.9)
    }

    # Bias towards certain patterns (simulate learning)
    if "carefully" in candidate.instruction_prefix.lower():
        metrics['relevance'] += 0.05
    if "structured" in candidate.output_format.lower():
        metrics['coherence'] += 0.05
    if "concise" in candidate.reasoning_style.lower():
        metrics['conciseness'] += 0.05

    # Clip to [0, 1]
    metrics = {k: min(v, 1.0) for k, v in metrics.items()}

return metrics

def _visualize_optimization_trajectory(self):
    """Visualize how optimization progressed"""
    df = pd.DataFrame(self.optimization_history)

    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))

    # Score progression
    ax1.plot(df['iteration'], df['best_score'], marker='o',
              linewidth=2, markersize=8, label='Best Score', color='green')
    ax1.plot(df['iteration'], df['avg_score'], marker='s',
              linewidth=2, markersize=8, label='Avg Score', color='blue')

```

```

        ax1.set_xlabel('Iteration', fontsize=12)
        ax1.set_ylabel('Score', fontsize=12)
        ax1.set_title('GEPA Optimization Trajectory', fontsize=14, fontweight='bold')
        ax1.legend()
        ax1.grid(True, alpha=0.3)

    # Pareto front size
    ax2.bar(df['iteration'], df['pareto_size'], color='coral', alpha=0.7)
    ax2.set_xlabel('Iteration', fontsize=12)
    ax2.set_ylabel('Pareto Front Size', fontsize=12)
    ax2.set_title('Pareto Front Evolution', fontsize=14, fontweight='bold')
    ax2.grid(True, alpha=0.3, axis='y')

    plt.tight_layout()
    plt.savefig('optimization_trajectory.png', dpi=300, bbox_inches='tight')
    plt.show()

def forward(self, document: str):
    """Generate summary using optimized prompt"""
    if self.optimized_prompt is None:
        raise ValueError("Must run optimize_prompts() first!")

    # Build optimized prompt
    full_prompt = self.optimized_prompt.prompt_template.format(
        prefix=self.optimized_prompt.instruction_prefix,
        format=self.optimized_prompt.output_format,
        reasoning=self.optimized_prompt.reasoning_style
    )

    full_prompt += f"\n\nDocument:\n{document}"

    # Use ChainOfThought with optimized prompt
    # (In practice, inject optimized prompt into signature)
    summarize = dspy.ChainOfThought(SummarizePDF)
    result = summarize(document=document)

    return result

```

## EVALUATION METRICS

```

In [15]: class SummaryEvaluator:
    """Comprehensive evaluation of summaries"""

    def __init__(self):
        self.embedding_model = SentenceTransformer('all-MiniLM-L6-v2')

    def evaluate(self, summary: str, original_doc: str, key_points: List[str]):
        """Compute multiple evaluation metrics"""

        metrics = {}

        # 1. Semantic similarity (relevance)
        doc_emb = self.embedding_model.encode([original_doc[:5000]])[0]

```

```

        summary_emb = self.embedding_model.encode([summary])[0]
        metrics['semantic_similarity'] = float(cosine_similarity(
            doc_emb.reshape(1, -1),
            summary_emb.reshape(1, -1)
        )[0][0])

    # 2. Compression ratio
    metrics['compression_ratio'] = len(summary) / len(original_doc)

    # 3. Key points coverage (heuristic)
    coverage = sum(1 for kp in key_points if any(
        word.lower() in summary.lower()
        for word in kp.split()[:3] # Check first 3 words
    )) / max(len(key_points), 1)
    metrics['key_points_coverage'] = coverage

    # 4. Coherence (sentence count heuristic)
    sentences = summary.split('.')
    metrics['coherence_score'] = min(len([s for s in sentences if len(s.st

    # 5. Overall quality score
    metrics['overall_score'] = np.mean([
        metrics['semantic_similarity'],
        1 - min(metrics['compression_ratio'], 1.0), # Prefer more compressed
        metrics['key_points_coverage'],
        metrics['coherence_score']
    ])

    return metrics

```

## COMPARATIVE ANALYSIS & VISUALIZATION

```

In [16]: def compare_summarizers(baseline_results: Dict, gepa_results: Dict,
                           save_path: str = 'comparison_results.png'):
    """Compare baseline vs GEPA results with visualizations"""

    # Prepare data
    metrics = list(baseline_results.keys())
    baseline_vals = list(baseline_results.values())
    gepa_vals = list(gepa_results.values())

    # Create comprehensive comparison plot
    fig = plt.figure(figsize=(16, 10))
    gs = fig.add_gridspec(3, 2, hspace=0.3, wspace=0.3)

    # 1. Bar chart comparison
    ax1 = fig.add_subplot(gs[0, :])
    x = np.arange(len(metrics))
    width = 0.35

    bars1 = ax1.bar(x - width/2, baseline_vals, width, label='Baseline DSPy',
                    color='skyblue', edgecolor='black', linewidth=1.5)
    bars2 = ax1.bar(x + width/2, gepa_vals, width, label='GEPA-Optimized',

```

```

        color='orange', edgecolor='black', linewidth=1.5)

ax1.set_xlabel('Metrics', fontsize=12, fontweight='bold')
ax1.set_ylabel('Score', fontsize=12, fontweight='bold')
ax1.set_title('Baseline vs GEPA: Metric Comparison', fontsize=14, fontweight='bold')
ax1.set_xticks(x)
ax1.set_xticklabels([m.replace('_', ' ') .title() for m in metrics], rotation=45)
ax1.legend(fontsize=11)
ax1.grid(True, alpha=0.3, axis='y')

# Add value

```

Comparision of summary side by side

```

In [17]: def visualize_gepa_side_by_side(gepa, baseline, width=50):
    """
    Safe visualization for GEPA vs Baseline.
    Handles missing attributes in baseline without crashing.
    """

    def get(obj, attr):
        return getattr(obj, attr, "") or ""

    sep = " | "
    line = "-" * (width * 2 + len(sep))

    def block(text):
        wrapped = textwrap.wrap(str(text), width=width)
        return wrapped if wrapped else [""]

    def format_two_columns(left, right):
        max_len = max(len(left), len(right))
        left += [""] * (max_len - len(left))
        right += [""] * (max_len - len(right))
        return "\n".join(f"{left[i]:<{width}}{sep}{right[i]:<{width}}" for i in range(max_len))

    print(line)
    print(f'{BASELINE}'.ljust(width)){sep}{'GEPA'.ljust(width)})"
    print(line)

    # Metadata
    baseline_meta = f"Confidence: {get(baseline, 'confidence')}\nLength: {get(baseline, 'length')}"
    gepa_meta      = f"Confidence: {get(gepa, 'confidence')}\nLength: {get(gepa, 'length')}"

    print(format_two_columns(block(baseline_meta), block(gepa_meta)))
    print(line)

    # Summary
    print(format_two_columns(block(get(baseline, "summary")), block(get(gepa, "summary"))))
    print(line)

    # Reasoning (baseline may not have this)
    print(format_two_columns(block(get(baseline, "reasoning")), block(get(gepa, "reasoning"))))
    print(line)

```

```

print(line)

# Key points
baseline_kp = "\n".join(get(baseline, "key_points").split("\n")) if get(baseline, "key_points") else ""
gepa_kp = "\n".join(get(gepa, "key_points").split("\n")) if get(gepa, "key_points") else ""

print(format_two_columns(block(baseline_kp), block(gepa_kp)))
print(line)

# Optional numeric comparison
comparison_df = pd.DataFrame({
    "Metric": ["Confidence", "Length (tokens)", "Coverage Score"],
    "Baseline": [
        get(baseline, "confidence"),
        get(baseline, "length_tokens"),
        get(baseline, "coverage_score")
    ],
    "GEPA": [
        get(gepa, "confidence"),
        get(gepa, "length_tokens"),
        get(gepa, "coverage_score")
    ]
})

return comparison_df

```

## Text Preprocessing and normalization

```

In [18]: def clean_text(text: str) -> str:
    if pd.isna(text):
        return ""
    text = str(text)
    text = text.lower()
    text = re.sub(r"http\S+|www\S+|https\S+", "", text) # remove urls
    text = re.sub(r"[^a-z0-9\s]", " ", text) # keep letters only
    text = re.sub(r"\s+", " ", text).strip()
    return text

def preprocess_text(text: str, remove_stopwords=True, lemmatize_tokens=True) -> str:
    text = clean_text(text)
    tokens = nltk.word_tokenize(text)
    if remove_stopwords:
        tokens = [t for t in tokens if t not in STOPWORDS]
    if lemmatize_tokens:
        tokens = [lemmatizer.lemmatize(t) for t in tokens]
    return " ".join(tokens)

```

```

In [20]: # 1. Extract PDF and preprocess
pdf_text = extract_pdf_text("/mnt/a/Projects/Prompt_Optimizer/pdfs/Book-Summarization.pdf")
normalized_text = preprocess_text(pdf_text, remove_stopwords=False)
# 2. Baseline
baseline_summarizer = BaselinePDFSummarizer()
baseline_result = baseline_summarizer(document=normalized_text)

```

```

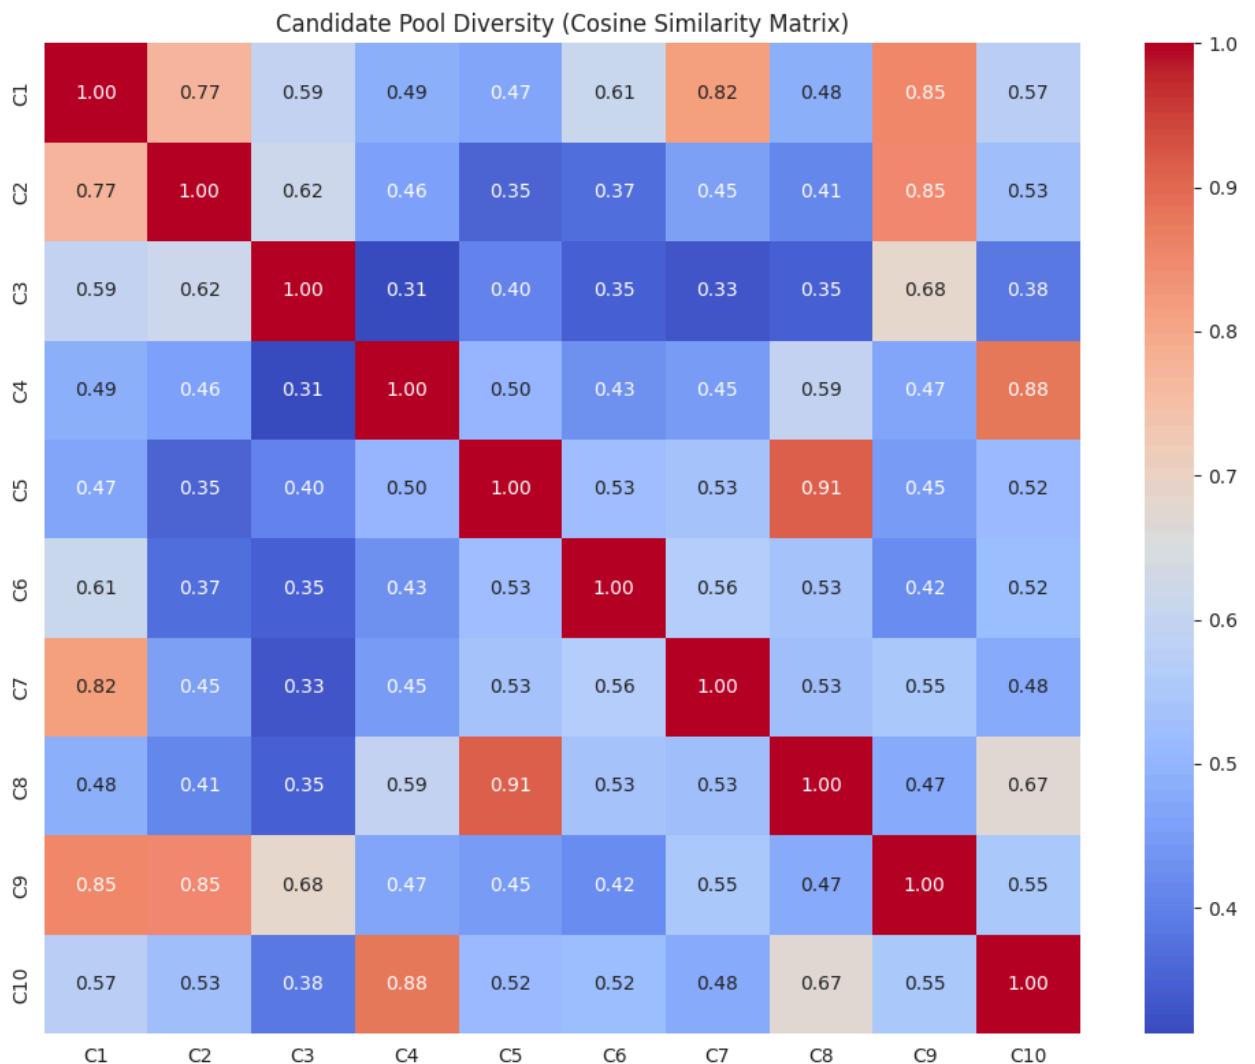
# 3. GEPA optimization
gepa_summarizer = GEPASummarizer(lm=lm, n_iterations=3)
gepa_summarizer.optimize_prompts([normalized_text]) # Run optimization
gepa_result = gepa_summarizer(document=normalized_text)

# 4. Evaluate & compare
evaluator = SummaryEvaluator()
baseline_metrics = evaluator.evaluate(baseline_result.summary, pdf_text, baseline_result.k)
gepa_metrics = evaluator.evaluate(gepa_result.summary, pdf_text, gepa_result.k)
visualize_gepa_side_by_side(gepa_result, baseline_result, width=50)
compare_summarizers(baseline_metrics, gepa_metrics)

```

=====
GEPA PROMPT OPTIMIZATION
=====

[Step 1] Initializing Candidate Pool...

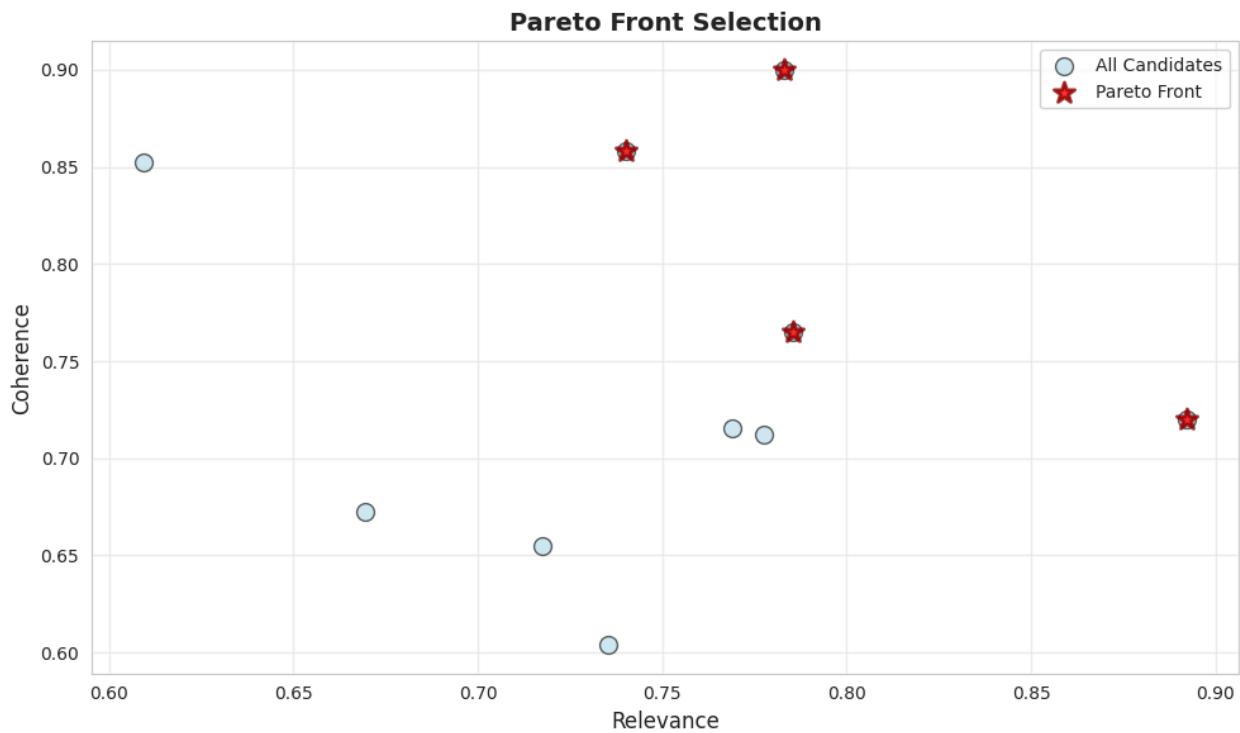


- ✓ Generated 10 diverse candidates  
Average pairwise similarity: 0.534

[Step 2] Evaluating Initial Candidates...

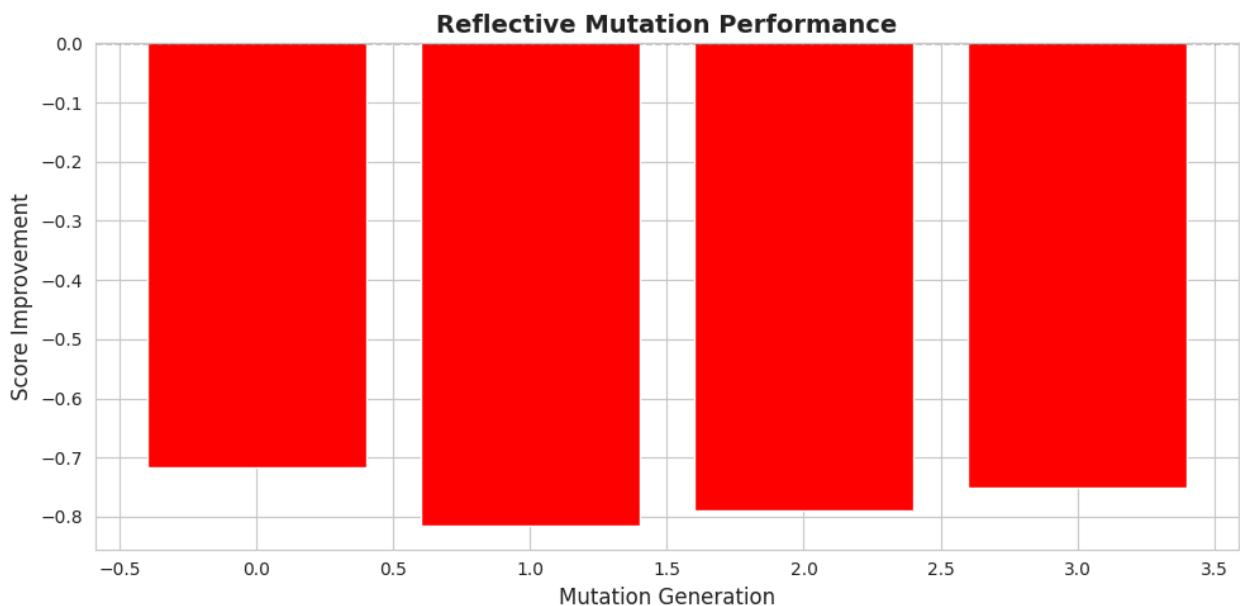
=====  
ITERATION 1/3  
=====

[Step 3.1] Applying Pareto Filtering...

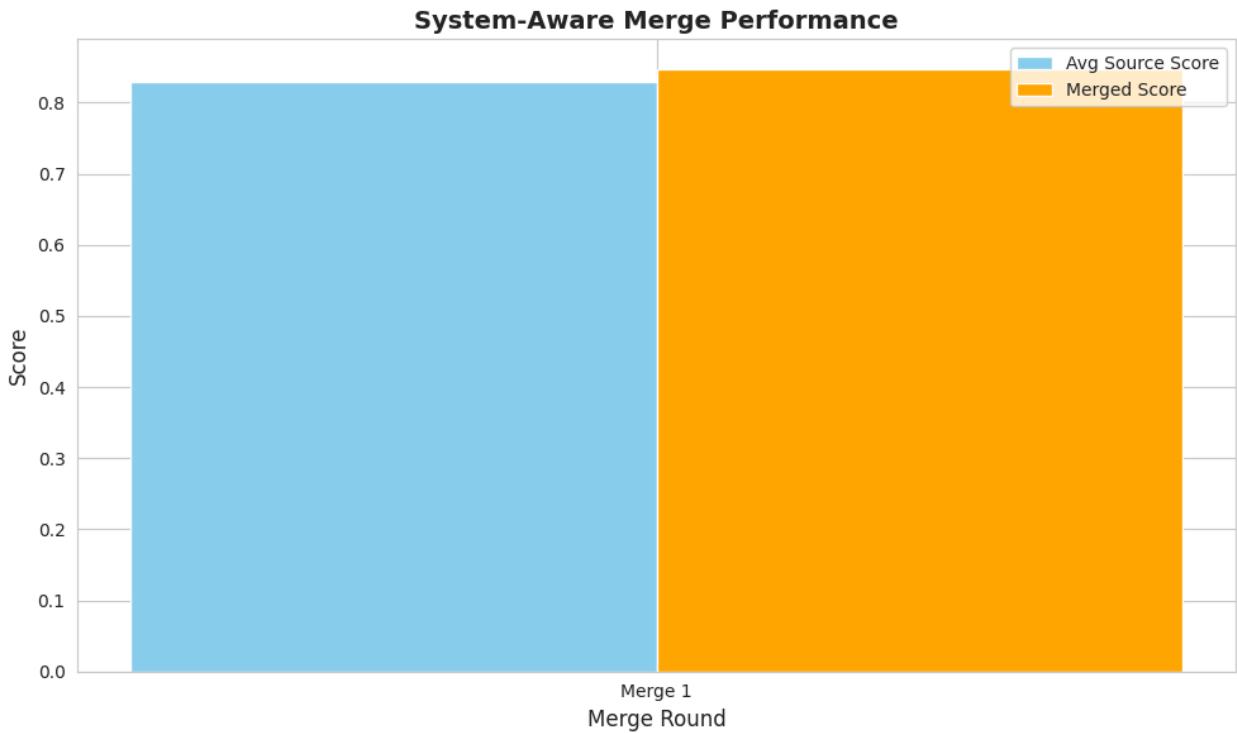


- ✓ Pareto front: 4/10 candidates

[Step 4.1] Applying Reflective Mutations...

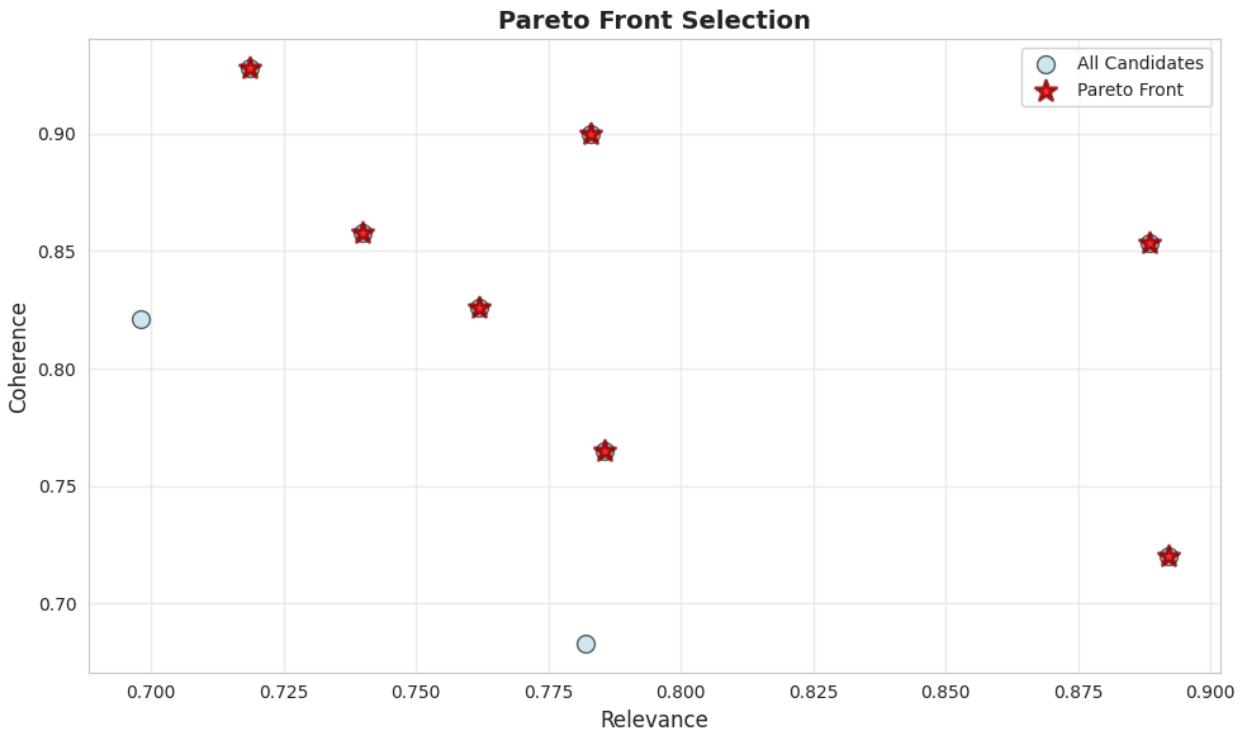


[Step 5.1] Merging Best Candidates...



=====  
ITERATION 2/3  
=====

[Step 3.2] Applying Pareto Filtering...

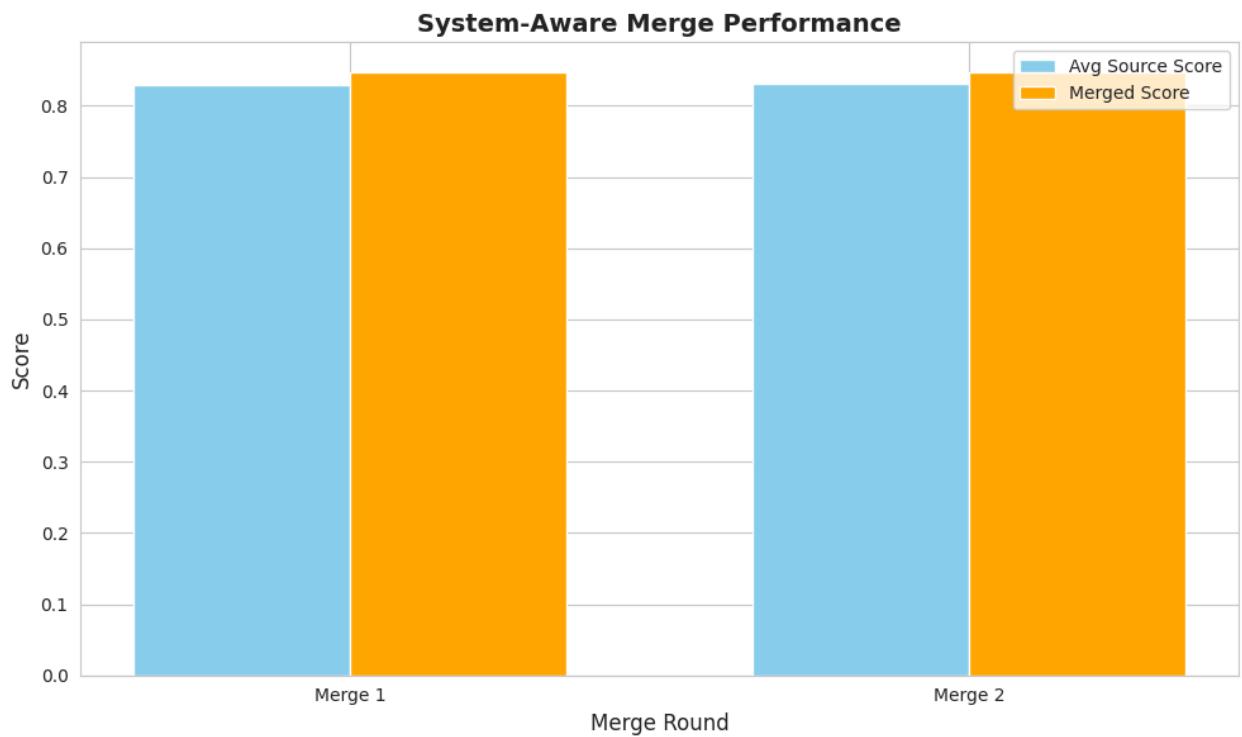


✓ Pareto front: 7/9 candidates

[Step 4.2] Applying Reflective Mutations...



[Step 5.2] Merging Best Candidates...

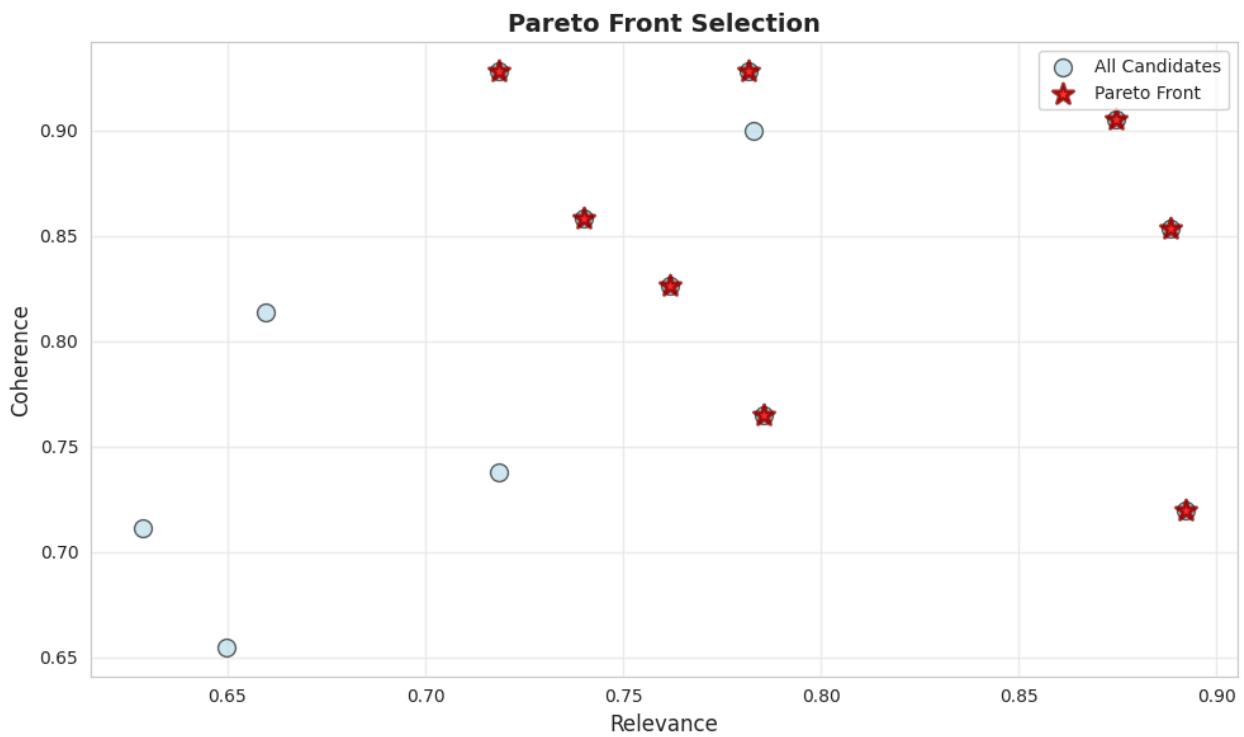


=====

ITERATION 3/3

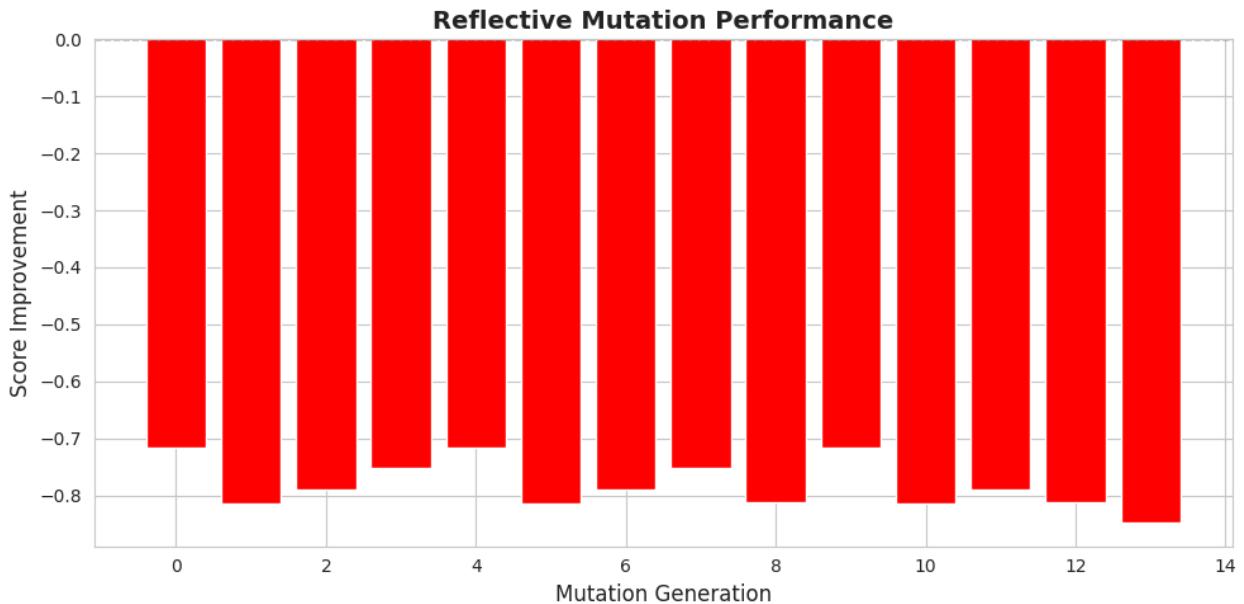
=====

[Step 3.3] Applying Pareto Filtering...



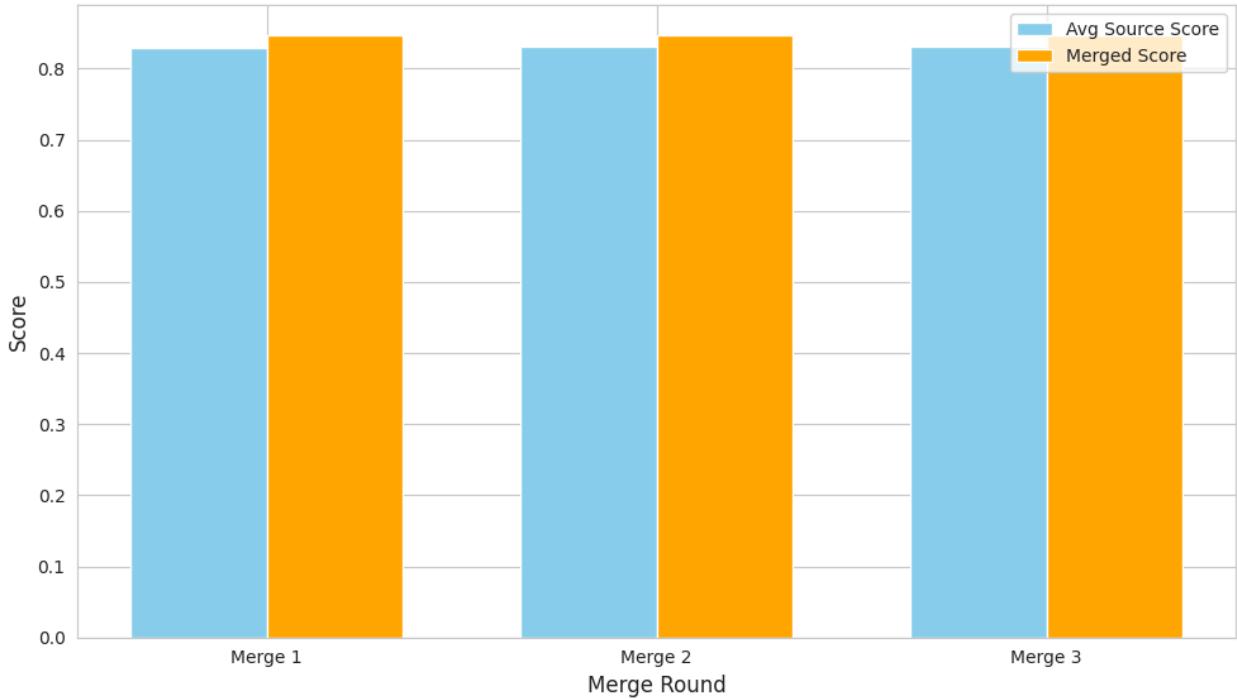
✓ Pareto front: 8/13 candidates

[Step 4.3] Applying Reflective Mutations...

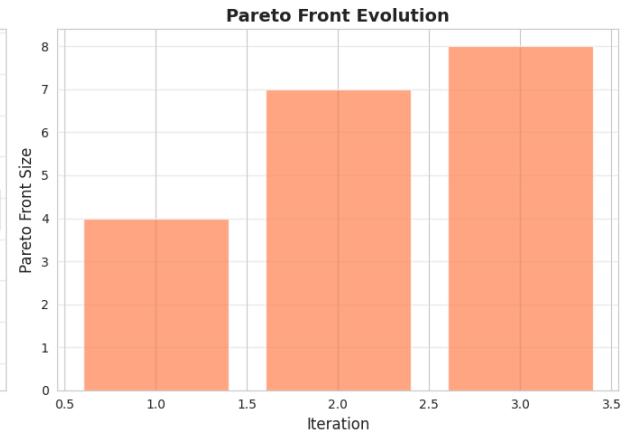
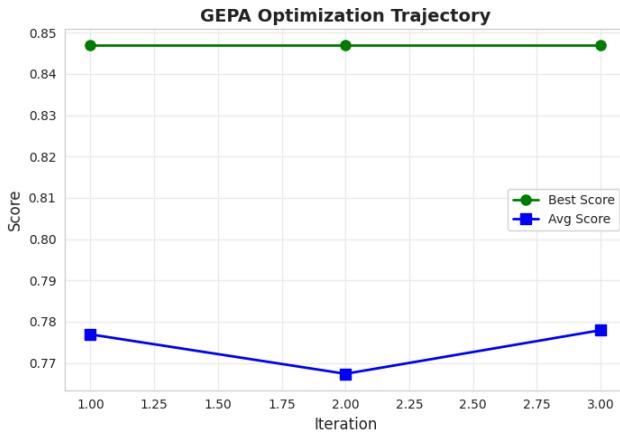


[Step 5.3] Merging Best Candidates...

### System-Aware Merge Performance



✓ Optimization Complete!  
Final Best Score: 0.847



-----  
-----  
BASELINE

GEPA
-----

Confidence: 0.8 Length: 256 Coverage: 0.9  
6 Coverage: 0.8

| Confidence: 0.9 Length: 25

The book "Rich Dad Poor Dad" by Robert Kiyosaki  
d" teaches readers how  
and Sharon Lechter teaches readers how to make  
om by having money work  
their money work for them. The author shares his  
s his personal story of  
personal experience of having two dads, one rich  
s, one who was poor and  
and one poor, and how they influenced his thoughts  
lessons he learned from  
on money. The book emphasizes the importance of  
the importance of  
financial literacy, minding one's own business,  
g your own business, and  
and using the cover of a corporation to achieve  
ration to achieve  
financial freedom. The author also discusses the  
hor also stresses the  
differences between the rich and the poor, and how  
ther than just working to  
the rich think and act differently when it comes  
to money. The book provides readers with practical  
advice on how to manage their finances, invest in  
assets, and build wealth.

| The book "Rich Dad Poor Da  
| to achieve financial freed  
| for them. The author share  
| growing up with two father  
| one who was rich, and the  
| them. The book emphasizes  
| financial literacy, mindin  
| using the cover of a corpo  
| financial freedom. The aut  
| need to work to learn, rat  
| earn a salary.  
|  
|

-----  
-----  
summary of the book  
ert Kiyosaki and Sharon  
es the importance of  
ing money work for you,  
around. The author  
ences and lessons  
rs, one who was  
ne who was not. The book  
chieving financial

| The document provided is a  
| "Rich Dad Poor Dad" by Rob  
| Lechter. The book emphasiz  
| financial literacy and hav  
| rather than the other way  
| shares his personal experi  
| learned from his two fathe  
| financially literate and o  
| highlights six rules for a  
| freedom, including not wor

king for money, having  
your own business,  
ration, inventing money,

| financial literacy, mindin  
| using the cover of a corpo  
| and working to learn.

-----  
-----  
\* The rich do not work for money, they have money  
money, but have money  
work for them \* Financial literacy is necessary to  
literacy is necessary to  
achieve financial freedom \* Minding one's own  
\* Minding your own  
business is crucial for building wealth \* Using  
r than just working for  
the cover of a corporation can provide tax  
of a corporation can  
benefits and protect assets \* Inventing money  
financial protection \*  
through creativity and innovation can lead to  
ng wealth, is a key  
financial success \* Working to learn, rather than  
rking to learn, rather  
just to earn, is essential for long-term financial  
is essential for  
growth \* The rich think and act differently when  
m \* Understanding the  
it comes to money, they prioritize financial  
and liabilities is  
education and take calculated risks  
e-generating assets is

| \* The rich do not work for  
| work for them \* Financial  
| achieve financial freedom  
| business is crucial, rathe  
| a salary \* Using the cover  
| provide tax benefits and f  
| Inventing money, or creati  
| principle of the book \* Wo  
| than just working to earn,  
| achieving financial freedo  
| difference between assets  
| critical \* Building revenu  
| key to achieving financial

freedom

